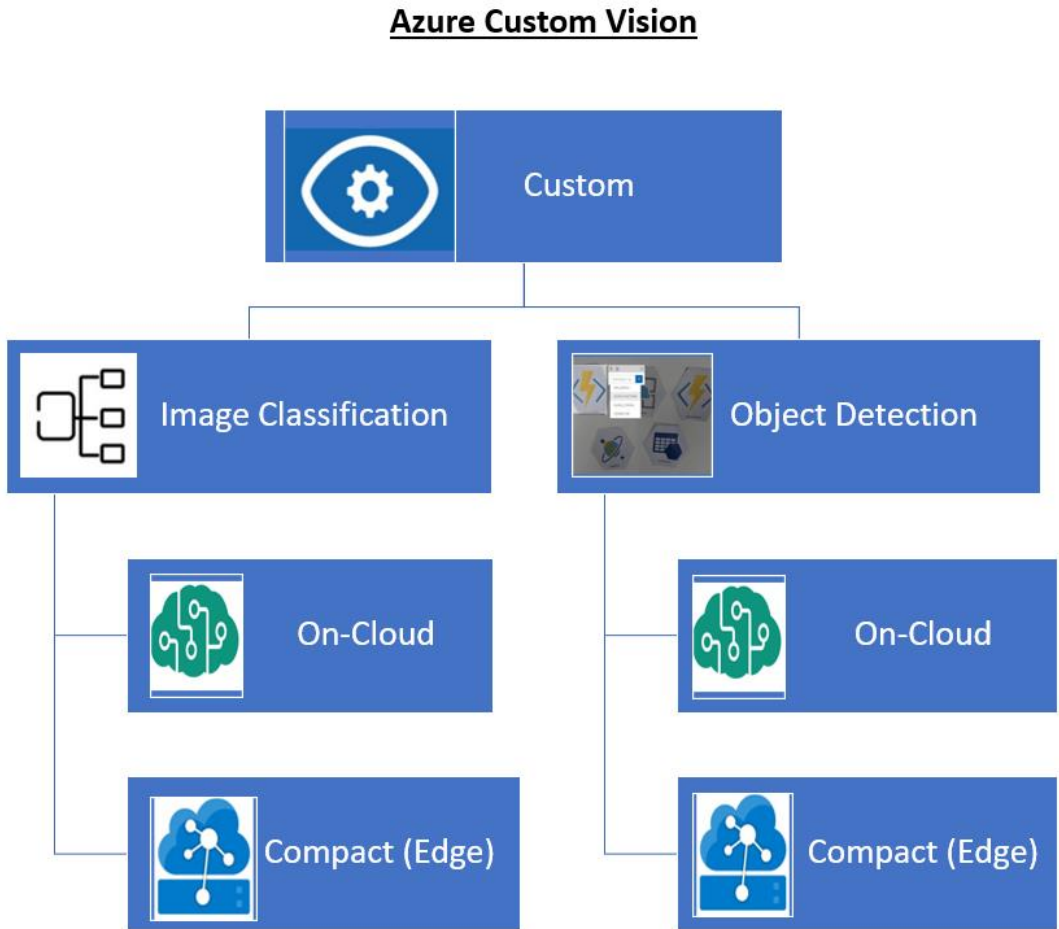


Document Name	HOL – Azure Custom Vision Service v3.x
Author	Shiva S Tomar & Anupreet Kaur
Reviewer	
Executive Summary	Azure Cognitive APIs enable the developers of all skill levels to add human intelligence in their applications. The services are designed for developers interested in pursuing DS/AI/ML skills and people who want to acquire the deep technical knowledge on the Cognitive APIs of Azure, despite not having Machine Learning expertise.
Purpose	<p>This document is created to help you gain level 350 working knowledge on Azure Custom Vision Service. You will be able to explore each functionality offered by the service through the GUI Portal to train the model & REST APIs to observe the outcomes. We have also shared a sample dataset to replicate what we have used to create the content of this workshop.</p> <p>Once you complete these labs, you’ll go from <i>Zero to Hero</i> on the respective Azure Cognitive service and should be able to <i>Demo, Develop and Deploy</i> your own custom use cases. The important thing to note here is that you don’t need to refer any other documents to complete this workshop.</p>
Intent of Guide	This workshop is designed to help you explore all the features of a service offered through their GUI Portal & APIs. The diagram shown in the beginning of the document is its functional Architecture; talking about the functionalities offered by the service in a flow. It also covers the Concepts, How-to and best practices about the service. This document is not intended to enable you with scenarios of deployment in production.

Service brief: Azure Custom Vision Service

Azure's Custom Vision service gives you access to build custom models for Image Classification & Object Detection using a no code or code first approach.

Diagram: Functional Architecture



Azure offers a separate Computer Vision service, that allows you to use pre-built models for Image Analysis & Object Detection. In case, you have a specific use case / requirements that may or may not be met by pre-built models, you can use the Custom Vision Service to create a custom model for your business.

Custom Vision models can be broadly classified into 2 categories :

1. Image Classification : This enables you to solve the classification problem pertaining to images. Using image classification, one can classify the images into 1 or more buckets by tagging the images. There are 2 types of classification problems -
 - a. Multiclass Classification : Each image is associated with 1 tag
 - b. Multilabel Classification : Each image can be associated with multiple tags
2. Object Detection : This enables you to detect objects present in an image by providing the bounding coordinates where they are present.

You can create & deploy the custom models in a no code approach using the [GUI portal](#) or in a code first approach using language specific SDKs or REST APIs. This workshop combines portal & REST API approach, for model building, training, testing, tuning and prediction.

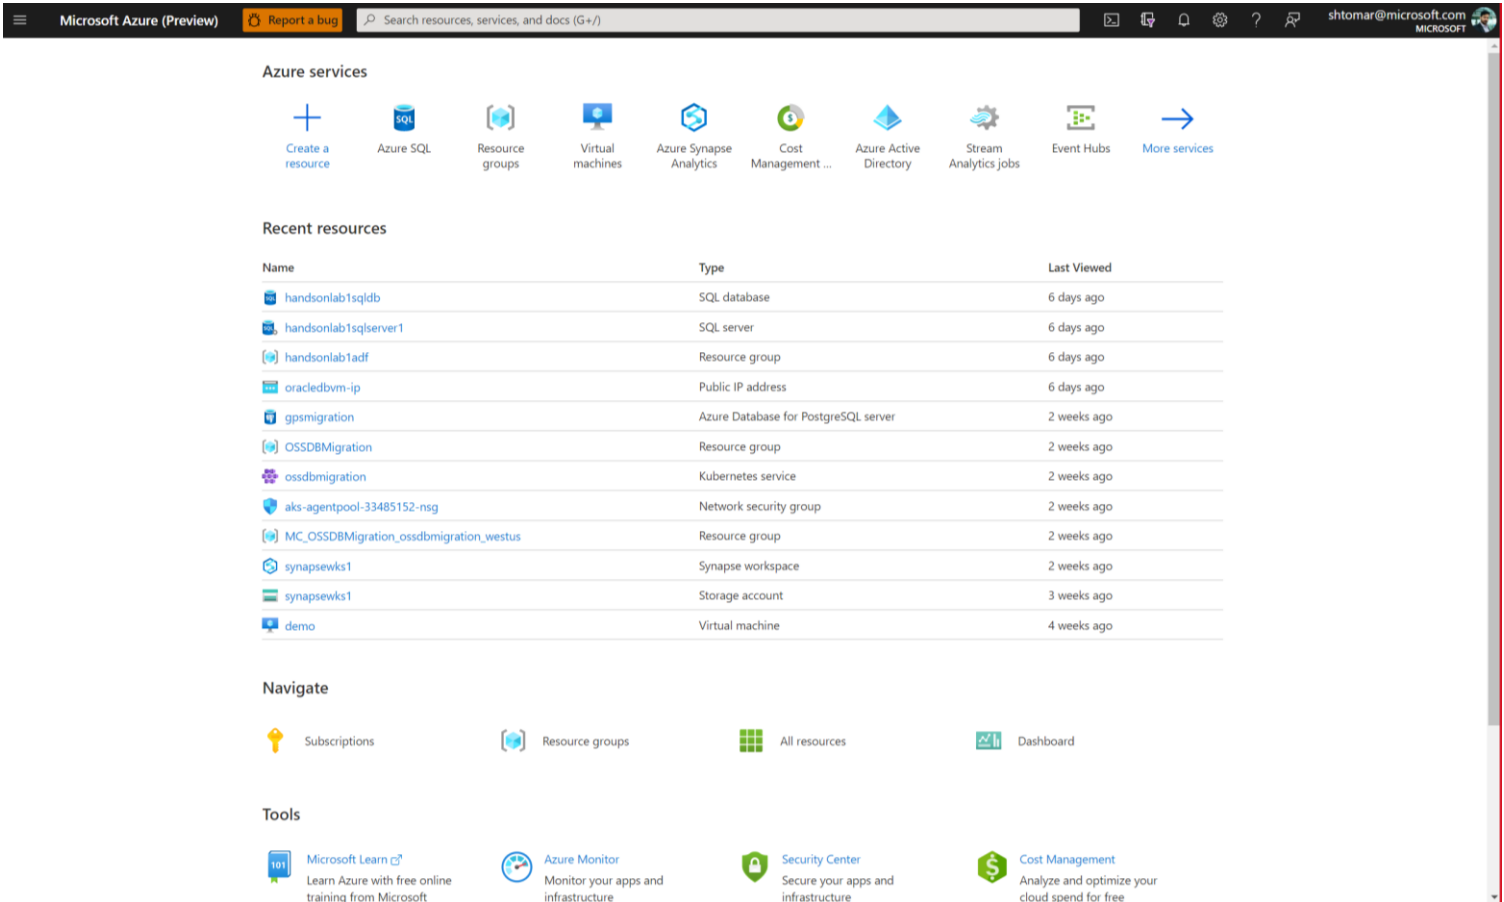
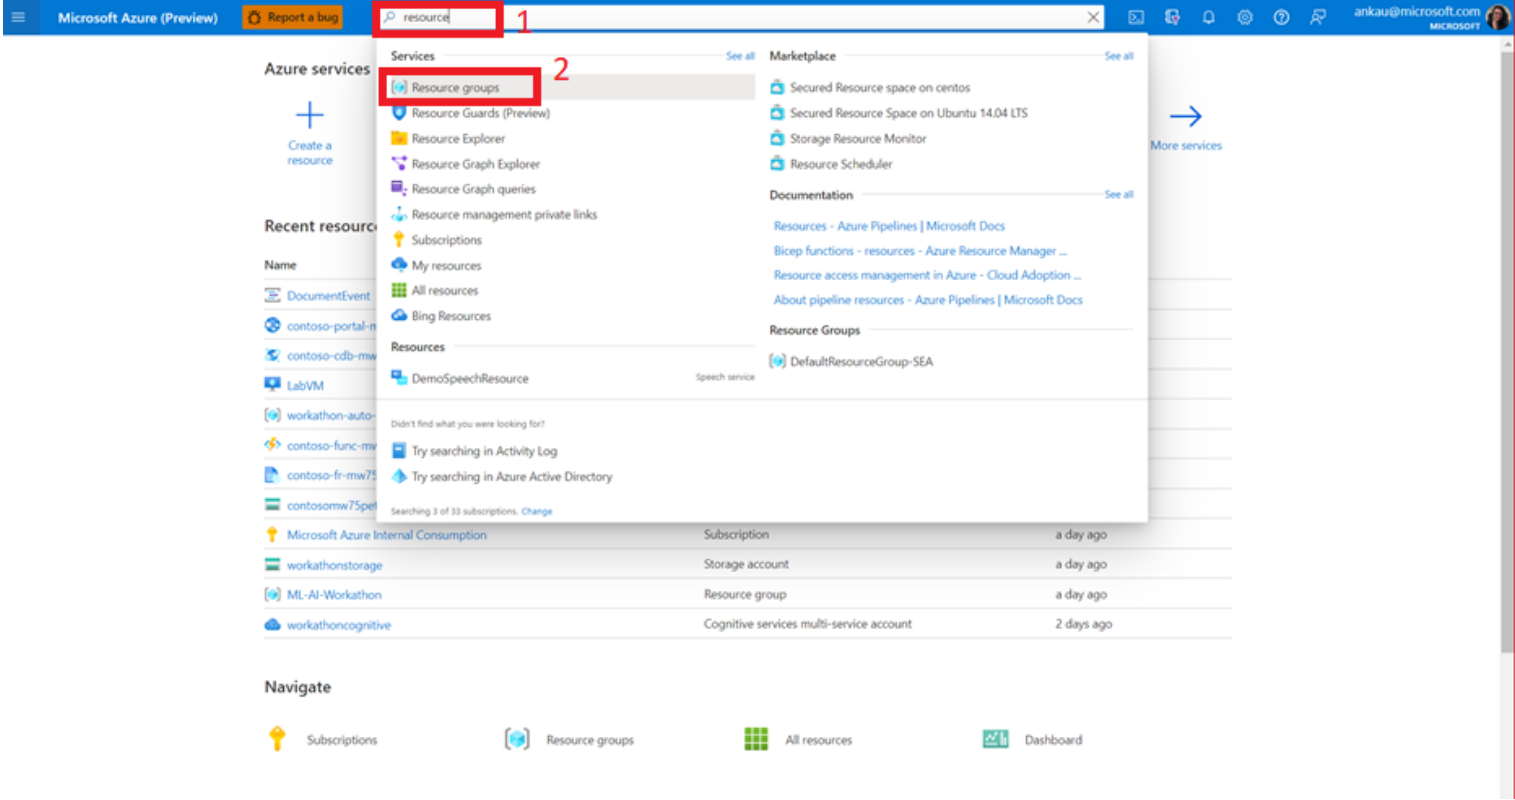
You can deploy the models on the cloud or on the edge.

Step by step hands on guide to go from *Zero to Hero*

Pre-requisites

- Download & Install Postman
 - Postman is a free tool which allows you to make API calls
 - You can download the desktop application or get started using the web version ([Download Postman | Try Postman for Free](#))
- An active Azure Account
 - You can use your current Azure Subscription or get started by creating a free trial account (<https://azure.microsoft.com/en-in/free>)
- Download the data for training & testing images from the Data folder in Custom Vision Folder

Let's get started!

Screenshots	Steps & Significance
	Sign into your Azure Portal.
	<p><u>Create a Resource Group</u></p> <p>Follow steps 1 & 2 to create a resource group.</p>

Microsoft Azure (Preview)

Report a bug

Search resources, services, and docs (G+/)

ankau@microsoft.com

Home >

Resource groups

Microsoft

+ Create

Manage view

Refresh

Export to CSV

Open query

Assign tags

Feedback

Filter for any field...

Subscription == 3 of 33 selected

Location == all

Add filter

Showing 1 to 24 of 24 records.

No grouping

List view

Name	Subscription	Location
<input type="checkbox"/> AzureMachineLearningDemo	anu_internalSubscription	Southeast Asia
<input type="checkbox"/> BotsDemo	anu_internalSubscription	Central India
<input type="checkbox"/> cloud-shell-storage-centralindia	anu_internalSubscription	Central India
<input type="checkbox"/> CognitiveServicesDemo	anu_internalSubscription	East US
<input type="checkbox"/> databricks-rg-DatabricksDemo1-wmlasg2Zkqe	anu_internalSubscription	Southeast Asia
<input type="checkbox"/> DatabricksMain	anu_internalSubscription	Southeast Asia
<input type="checkbox"/> DefaultResourceGroup-SEA	anu_internalSubscription	Southeast Asia
<input type="checkbox"/> formrecogdemo1	anu_internalSubscription	Central India
<input type="checkbox"/> FormRecognizerDemo	anu_internalSubscription	East US
<input type="checkbox"/> handsonlab1adff	anu_internalSubscription	Central India
<input type="checkbox"/> handsonlab1adff-synapse	anu_internalSubscription	Central India
<input type="checkbox"/> ML-AI-Workathon	Microsoft Azure Internal Consumption	Central India
<input type="checkbox"/> NetworkWatcherRG	anu_internalSubscription	East US
<input type="checkbox"/> ossdb	anu_internalSubscription	Central India
<input type="checkbox"/> QnAMakeDemo	anu_internalSubscription	West US
<input type="checkbox"/> rg0906	anu_internalSubscription	Southeast Asia
<input type="checkbox"/> synapserandom	anu_internalSubscription	Central India
<input type="checkbox"/> synapseworkshop	anu_internalSubscription	Central India

Click create to create a new resource group.

Microsoft Azure (Preview)

Report a bug

Search resources, services, and docs (G+/)

ankau@microsoft.com

Home > Resource groups >

Create a resource group

Microsoft

Basics

Tags

Review + create

Resource group - A container that holds related resources for an Azure solution. The resource group can include all the resources for the solution, or only those resources that you want to manage as a group. You decide how you want to allocate resources to resource groups based on what makes the most sense for your organization. Learn more

Project details

Subscription

anu_internalSubscription

Resource group

ML-AI-Workathon

Resource details

Region

(Asia Pacific) Central India

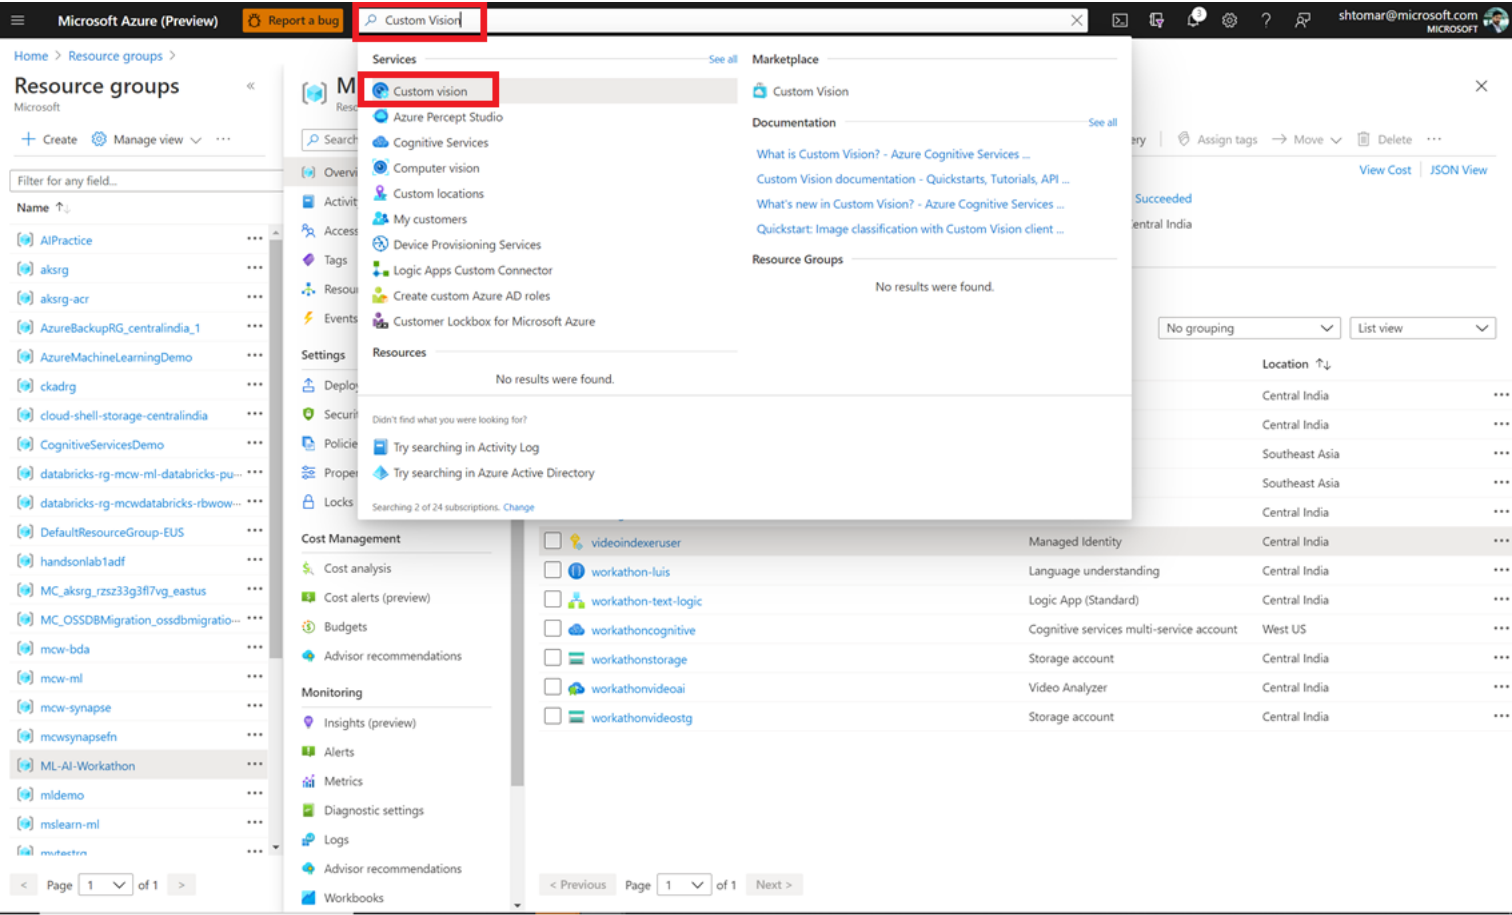
Review + create

< Previous

Next : Tags >

- Enter the details –
- Subscription : Azure subscription in which you want to deploy the resource group
 - Resource Group : Name of your choice for the resource group
 - Region : Region where you want to deploy the resource group

Click Review + Create.



Create a Custom Vision Service resource

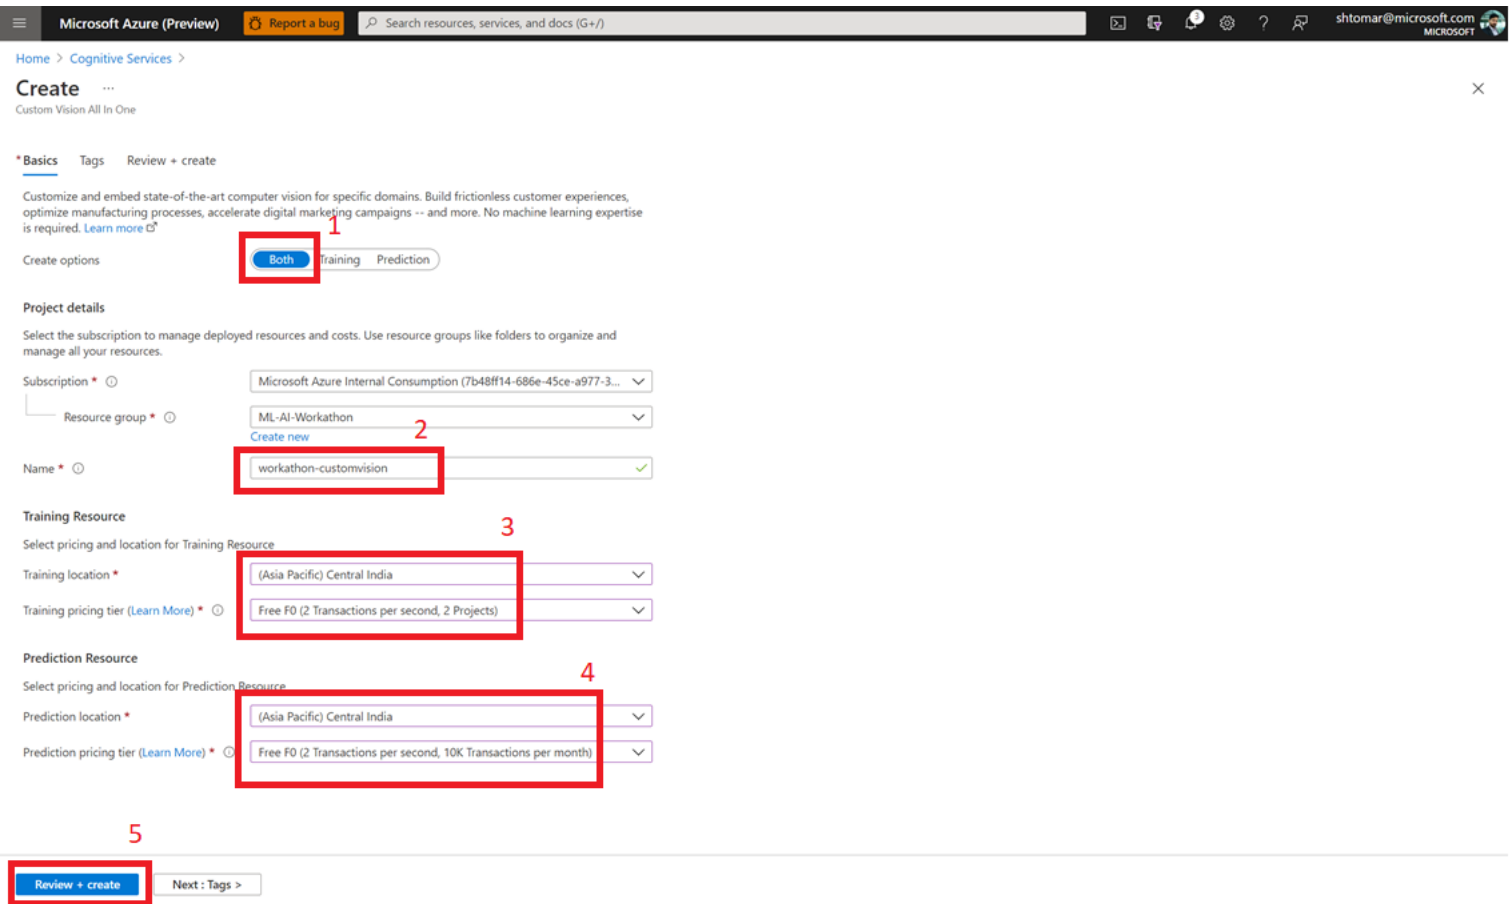
We are required to choose out of 3 resource type options on the Create custom service page, as shown in step 1 -

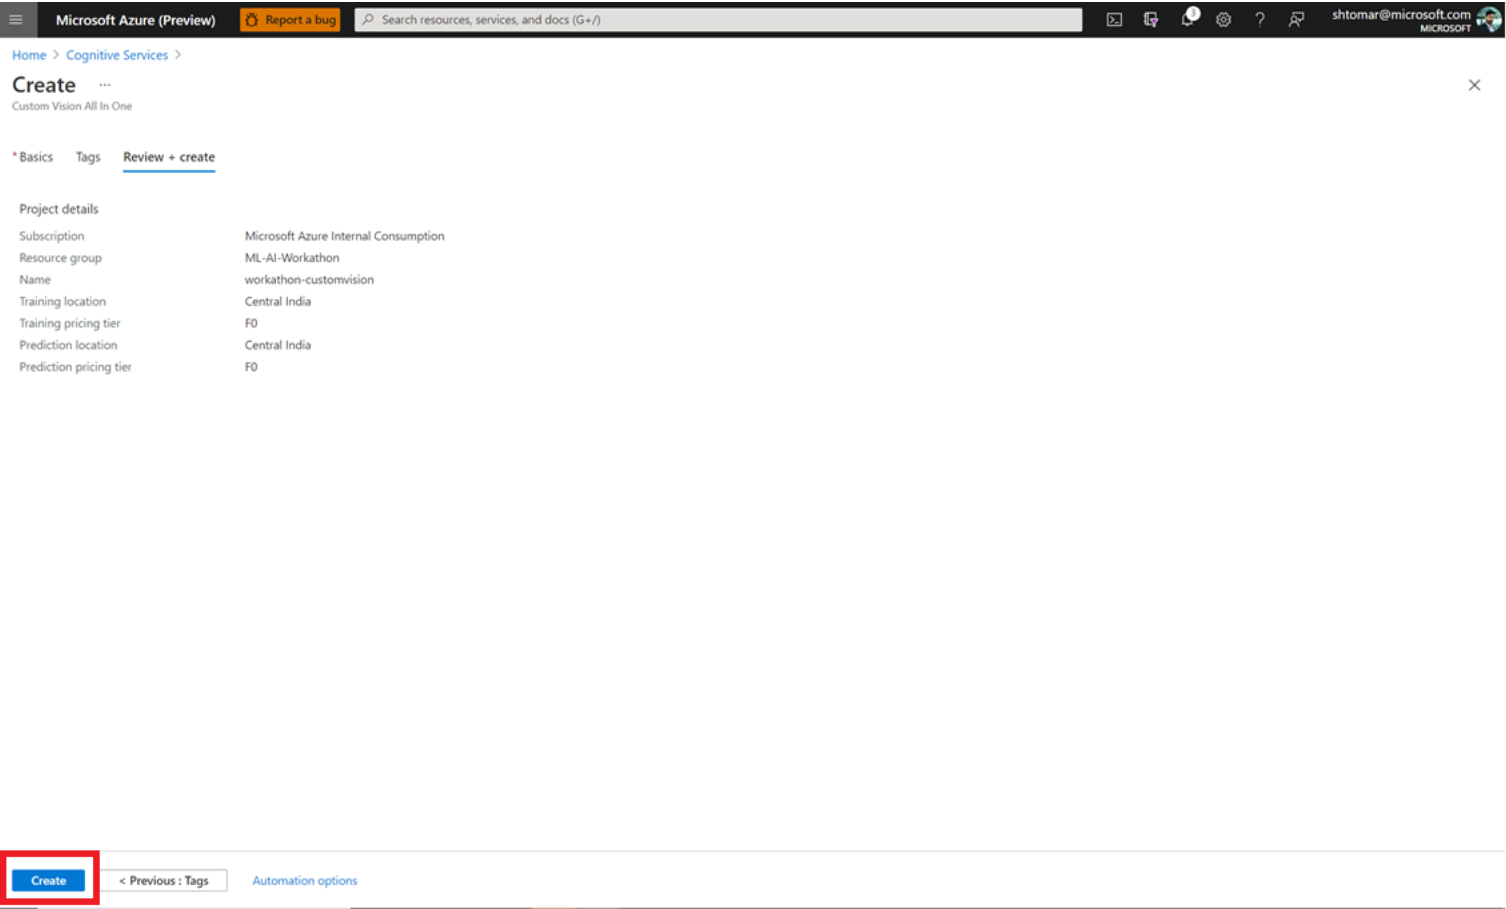
1. Training : Choosing this option will allow you to leverage this resource only to train a custom model. You will need a separate resource for Deployment & Prediction.
2. Prediction : Choosing this option will allow you to leverage this resource only for Deployment & Prediction of the trained model. You will need a separate resource to train a custom model.
3. Both : Choosing this option will allow you to spin up both resources at once, in the region of your choice. Hence, we are choosing this.

You might consider spinning resources separately in case you want to reuse a previously created resource for either Training or Prediction.

You could also train in 1 region & resource and choose to deploy to multiple Prediction resources in multiple regions, bases your use case.

Fill the details shown in step 2-4. Once done, click Review + Create.





Verify the details and click Create.

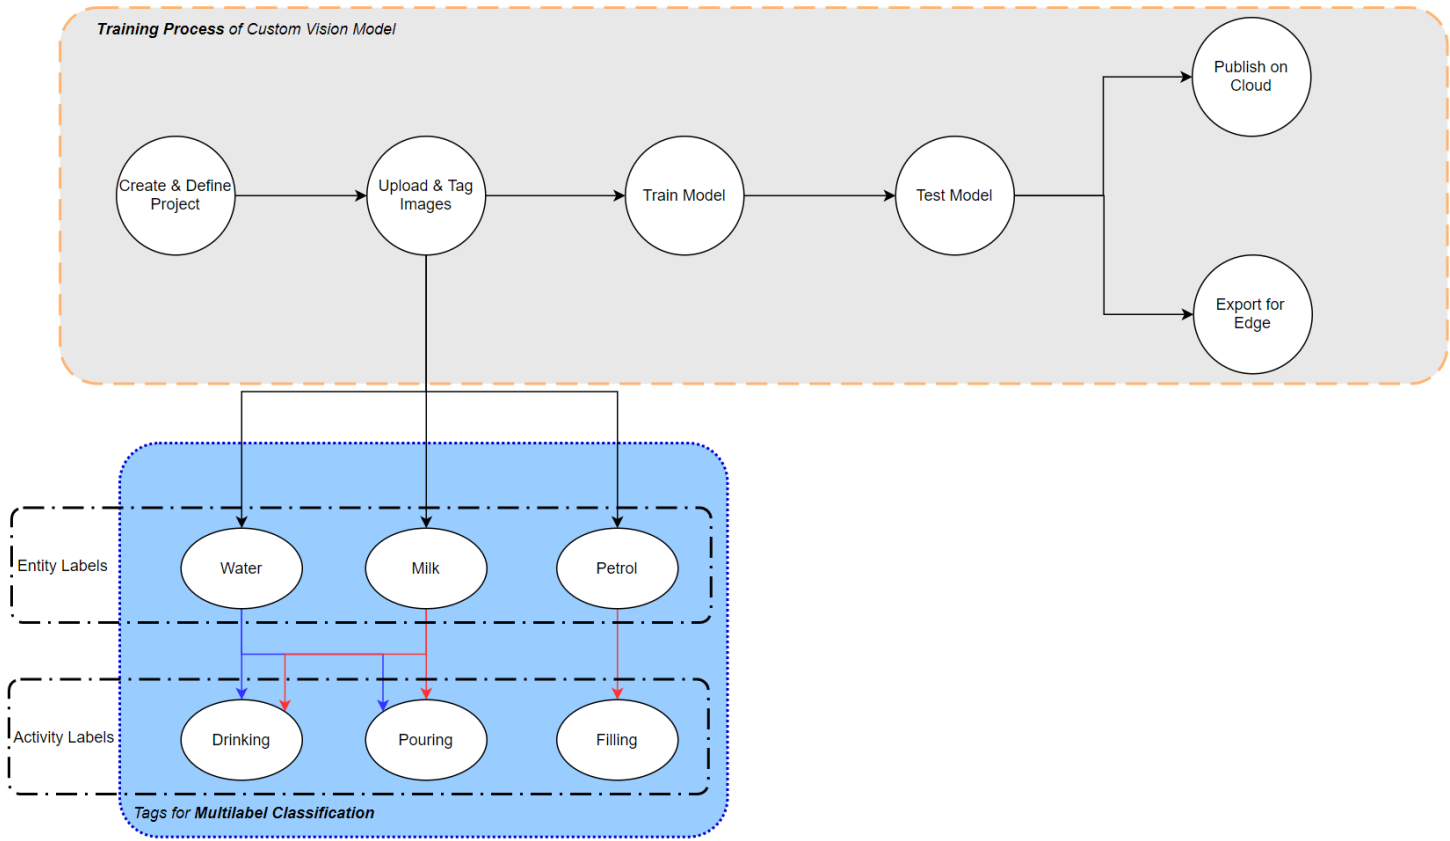
Once provisioned, go back to the Resource group.

You will see 2 resources :

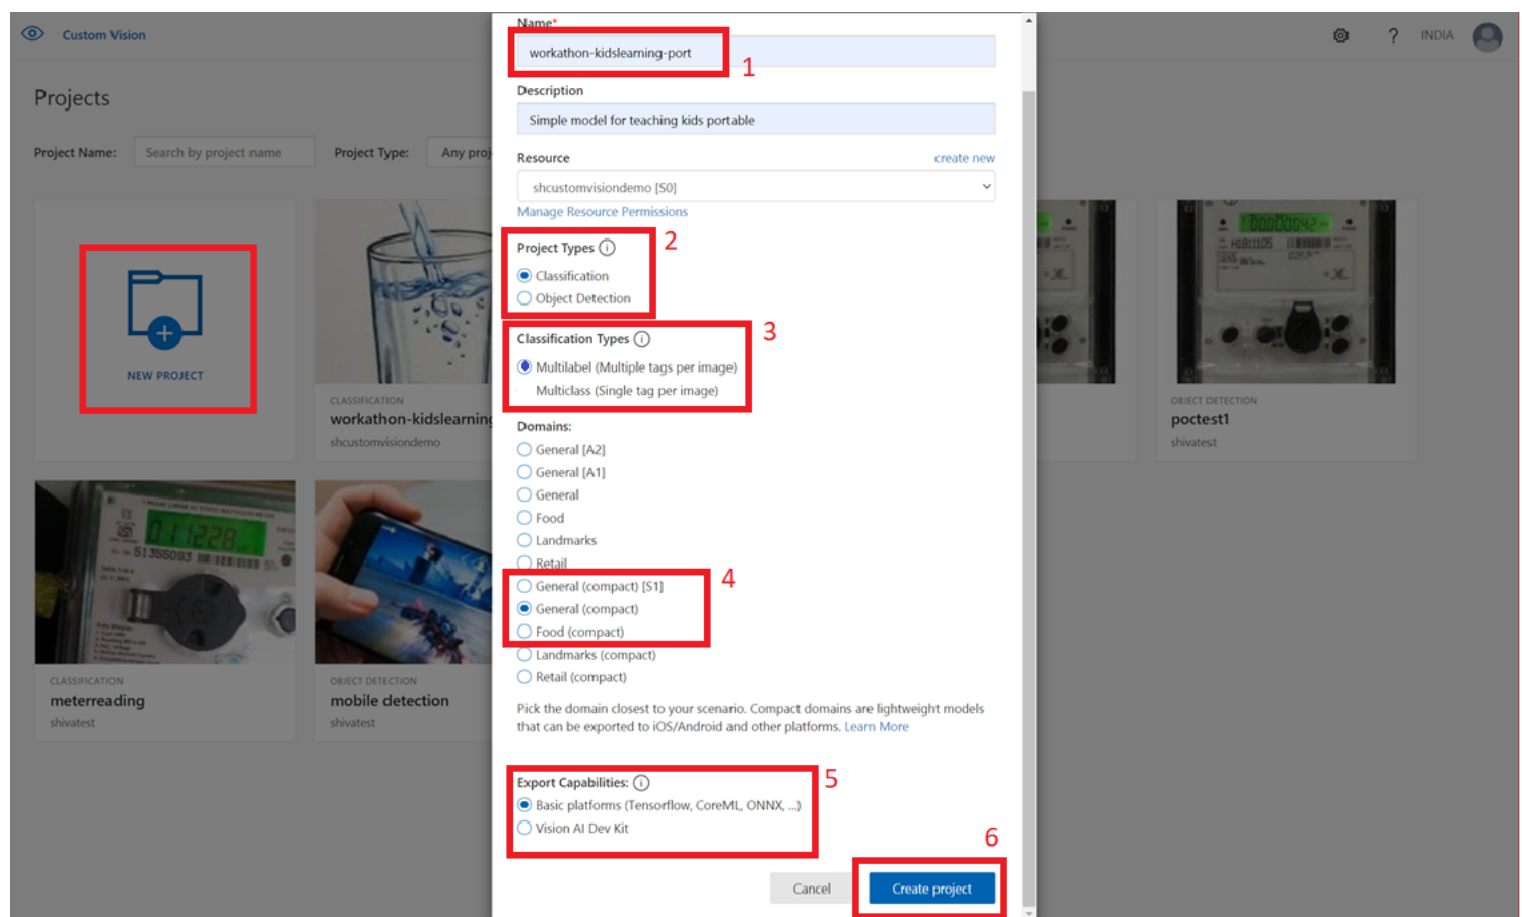
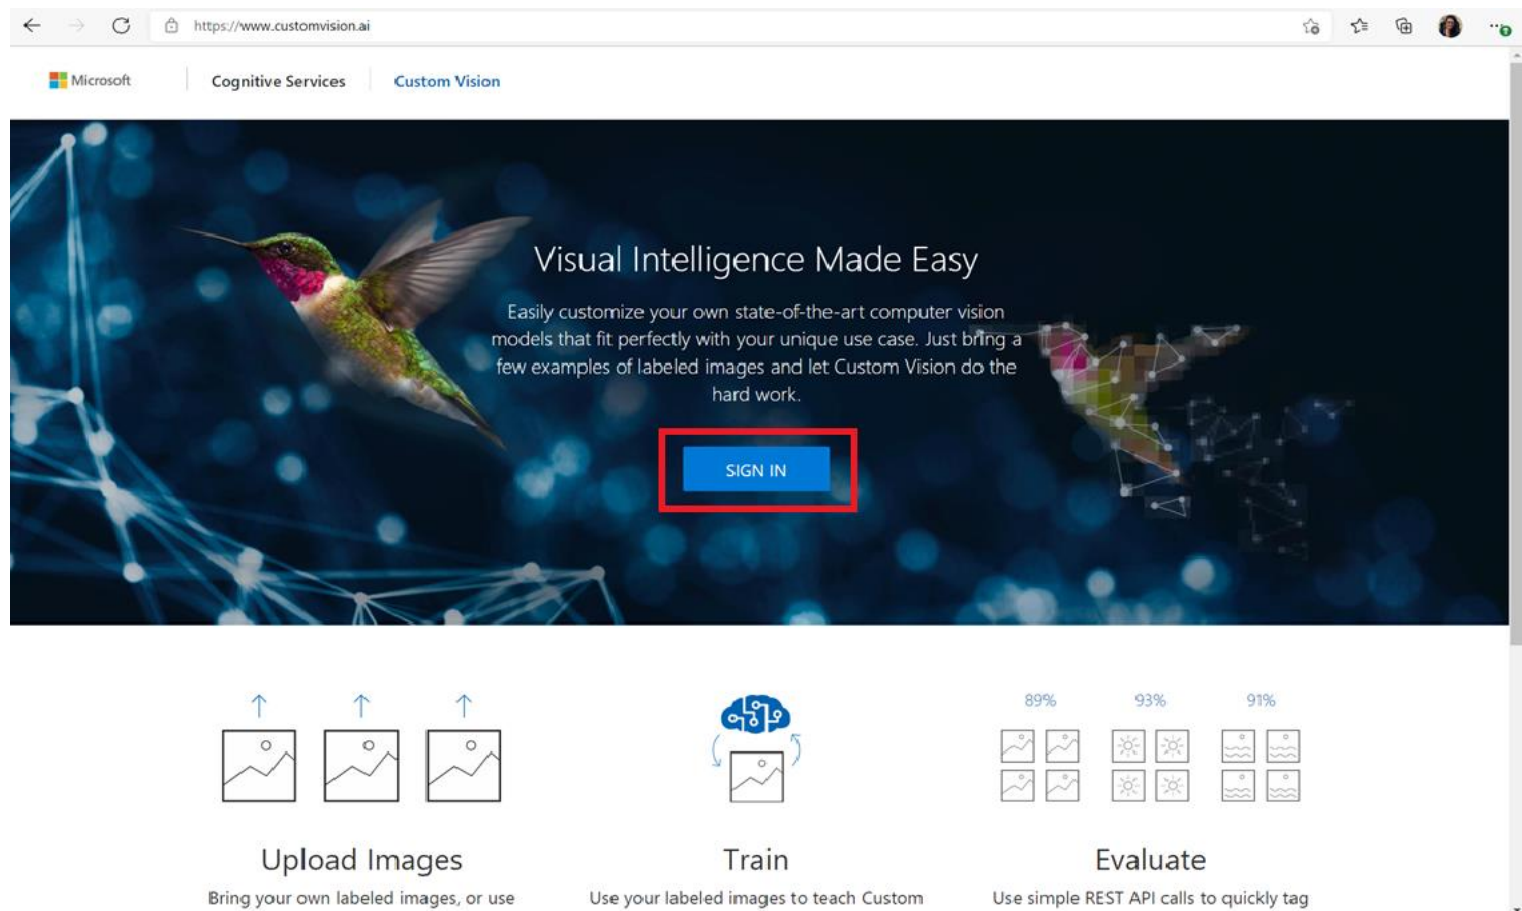
- 1. Workathon-customvision (This is the training resource)
- 2. Workathon-customvision-pred (This is the prediction resource)

Custom Model Training Workflow

In this workshop, you will be training a model using Custom Vision to perform multi label classification on any image. You will begin with creating & defining a project in Custom Vision Portal. Once you have created a project, you will be uploading the images & tagging them as Entities & Activities mentioned in the blue box below. In the diagram, Entities & Activities appear in a hierarchical format. However, this is just for your better understanding. When you create a multilabel project in Custom Vision, all tags will appear at the same level and not hierarchically. We will then train, test & publish / export the model for consumption in production.



Multilabel Classification simply means you can have multiple tags per image. For Eg : In an image with a person drinking water, we can tag that as Water (Entity) & Drinking (Activity). Multilabel classification is an extension to Multiclass Classification (single tag per image). Hence, performing this workshop will cover the concepts of both.

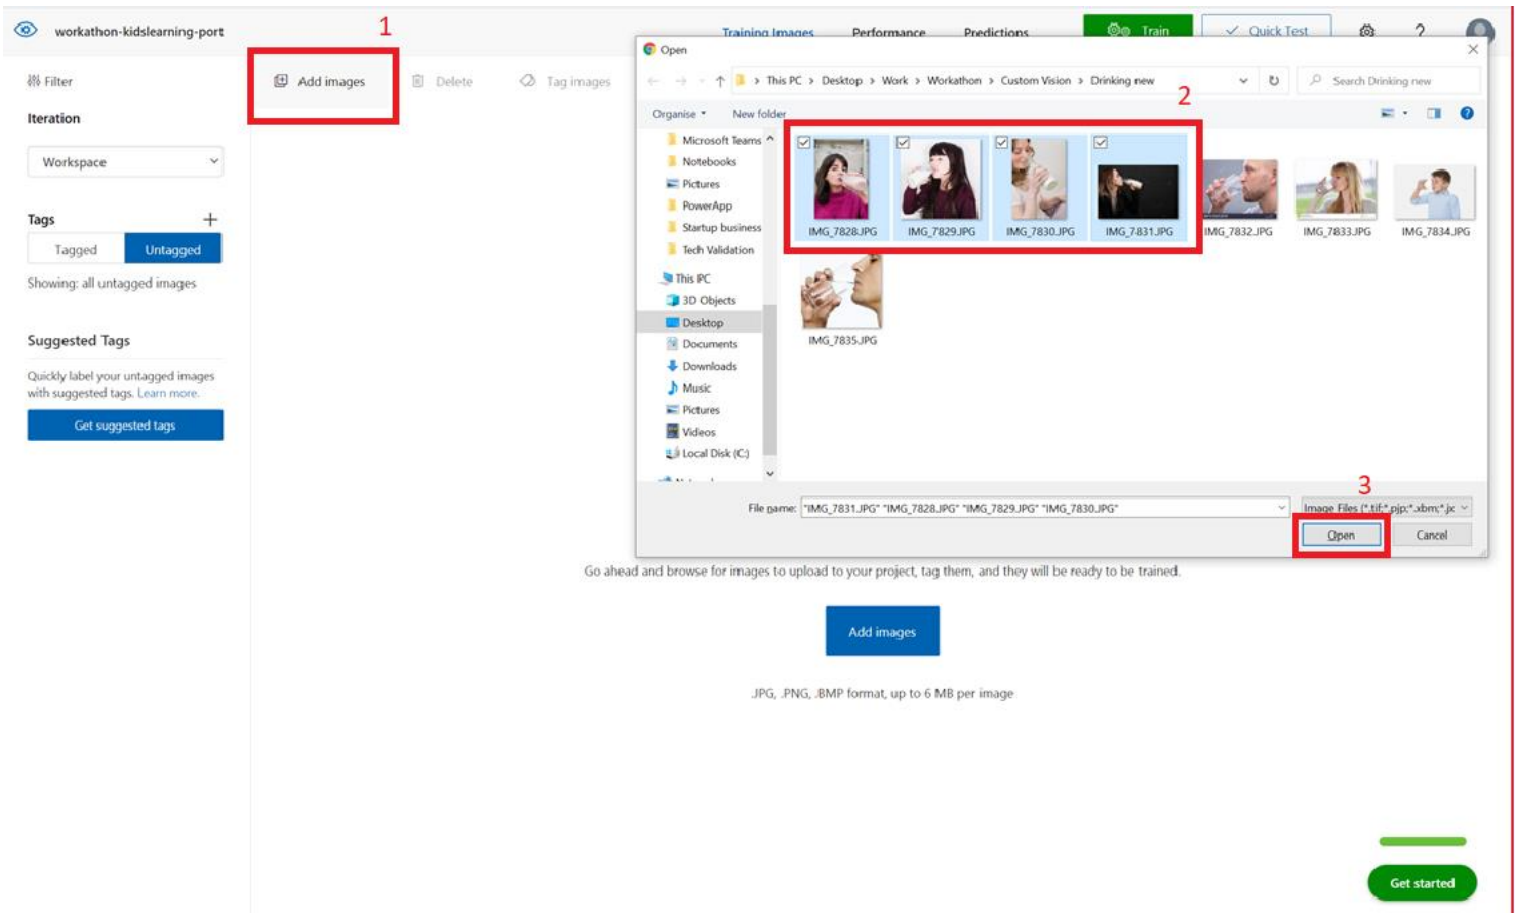


Create a Custom Vision Multilabel Classification Project

1. Open the [Custom Vision Portal](#)
2. Sign in with Azure account
3. Enter the project details –
 - a. Name of your choice
 - b. Project type : Classification
 - c. Classification type : Multiclass (Multiple tags per image)
 - d. Domain : General (compact)
 - e. Export Capabilities : Basic platforms
4. Click Create Project

Significance

1. Domains : General, Food, Retail, Landmarks etc. This allows you to select some pre-built models that lead to higher accuracy when trained with images from that domain. If your use case doesn't fit any of these domains, choose a General Domain.
2. Compact Models : Each of the available Domains have a corresponding compact model available as well. Selecting the compact models enables you to export these models to deploy them on edge. You can change to a compact domain later & retrain as well, however, it is recommended to select this before hand since there could be some difference in results between a compact vs non-compact domain.
3. Export capabilities : It allows you to choose a platform of your choice in which you would like to export.
Click [this](#) to learn more.



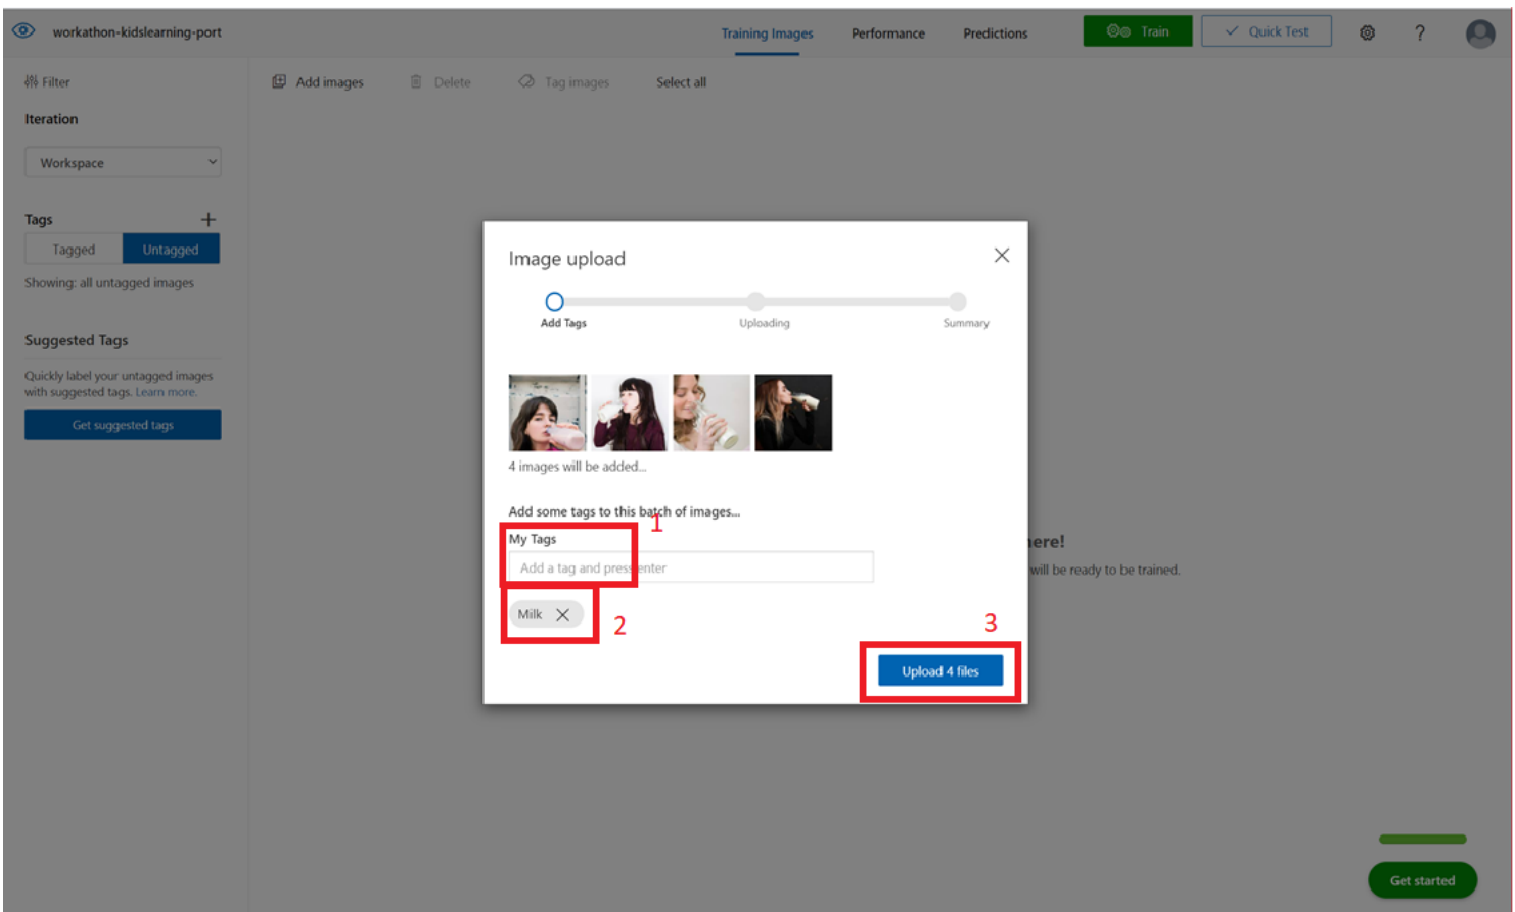
Adding Images for training

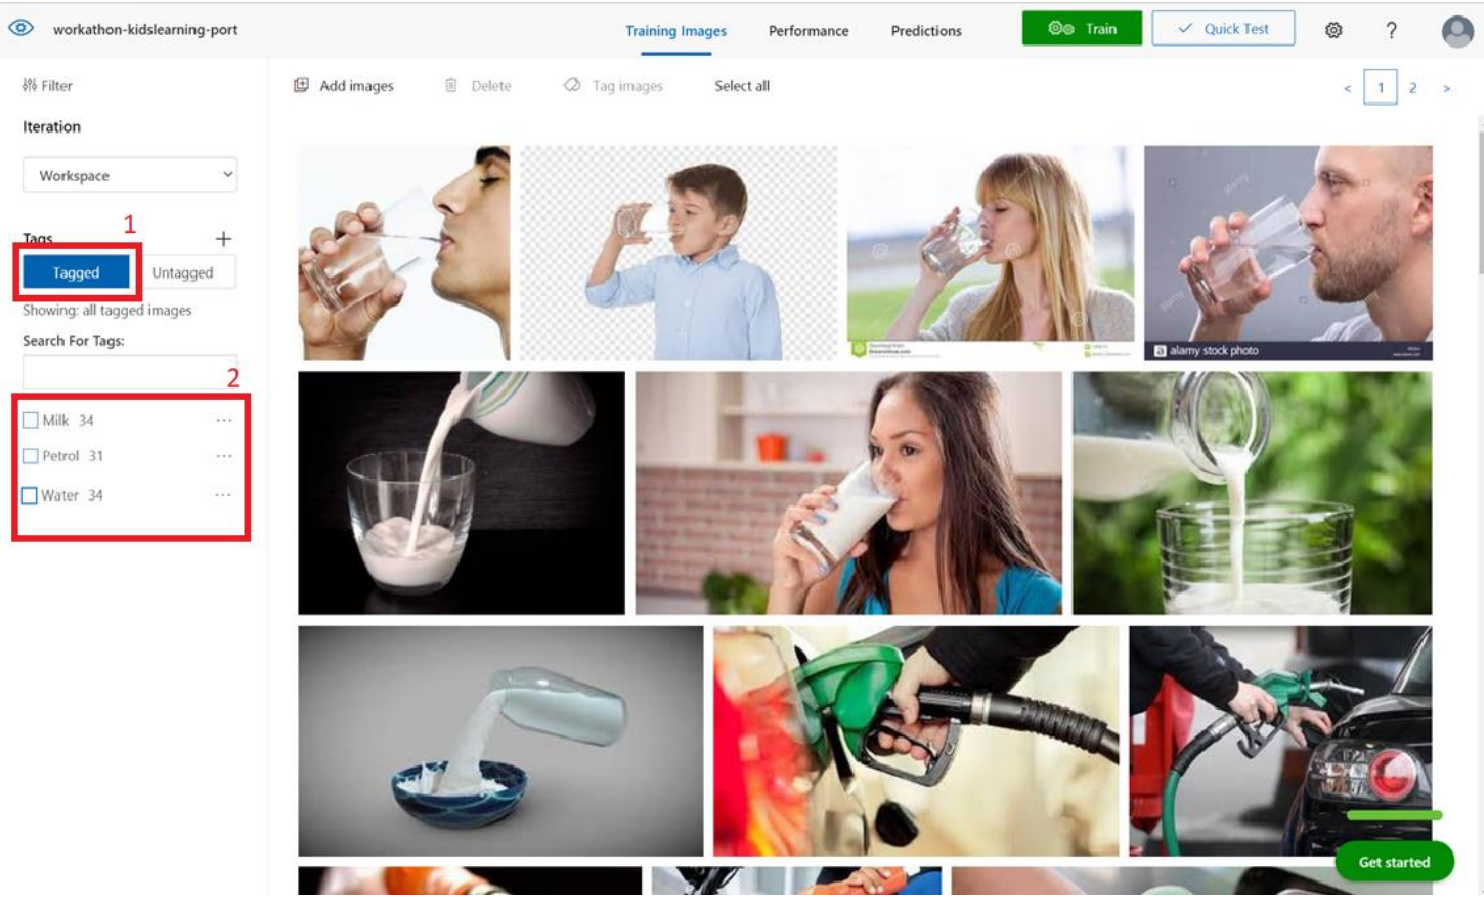
In this step, we will add images for training.
Follow the steps as highlighted.

We will first tag images as Milk, Water & Petrol.

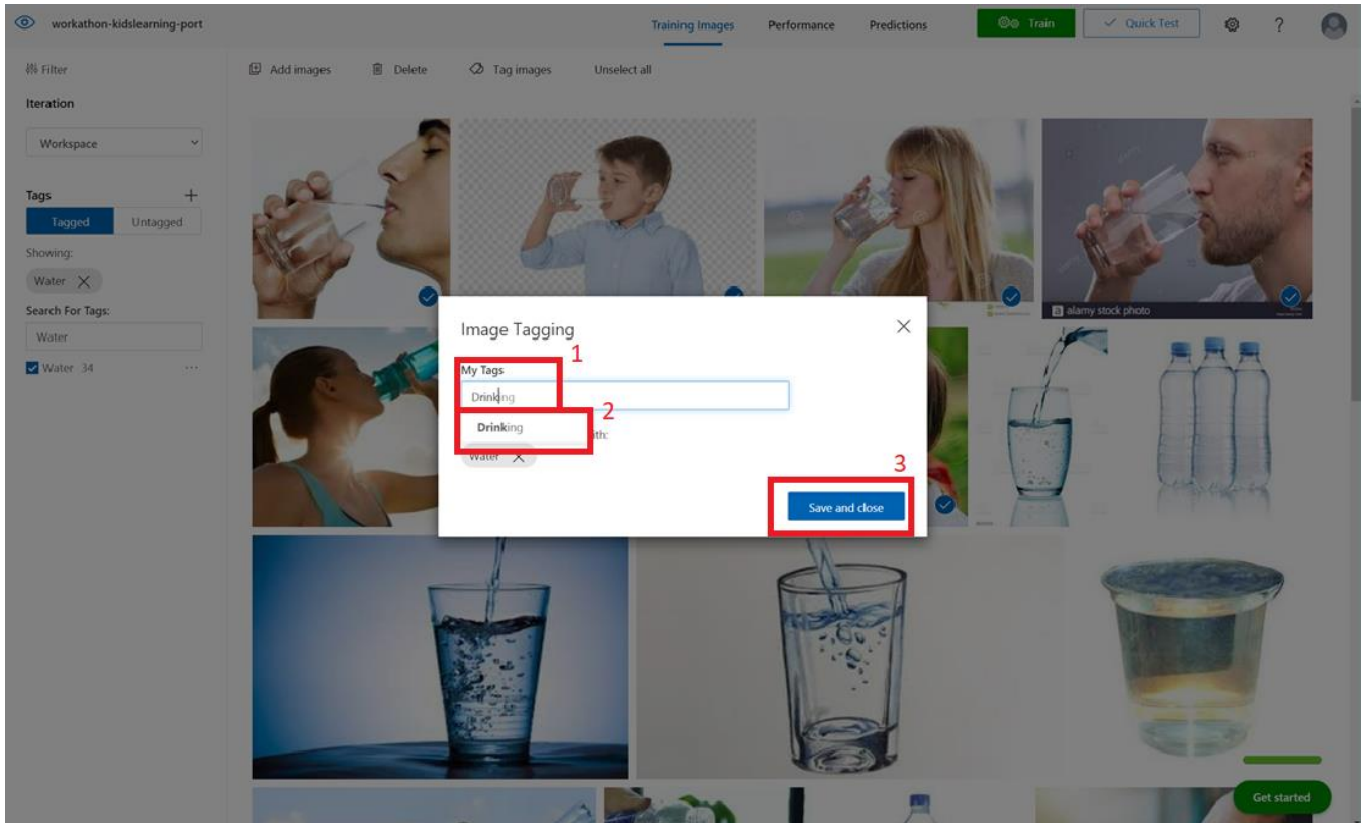
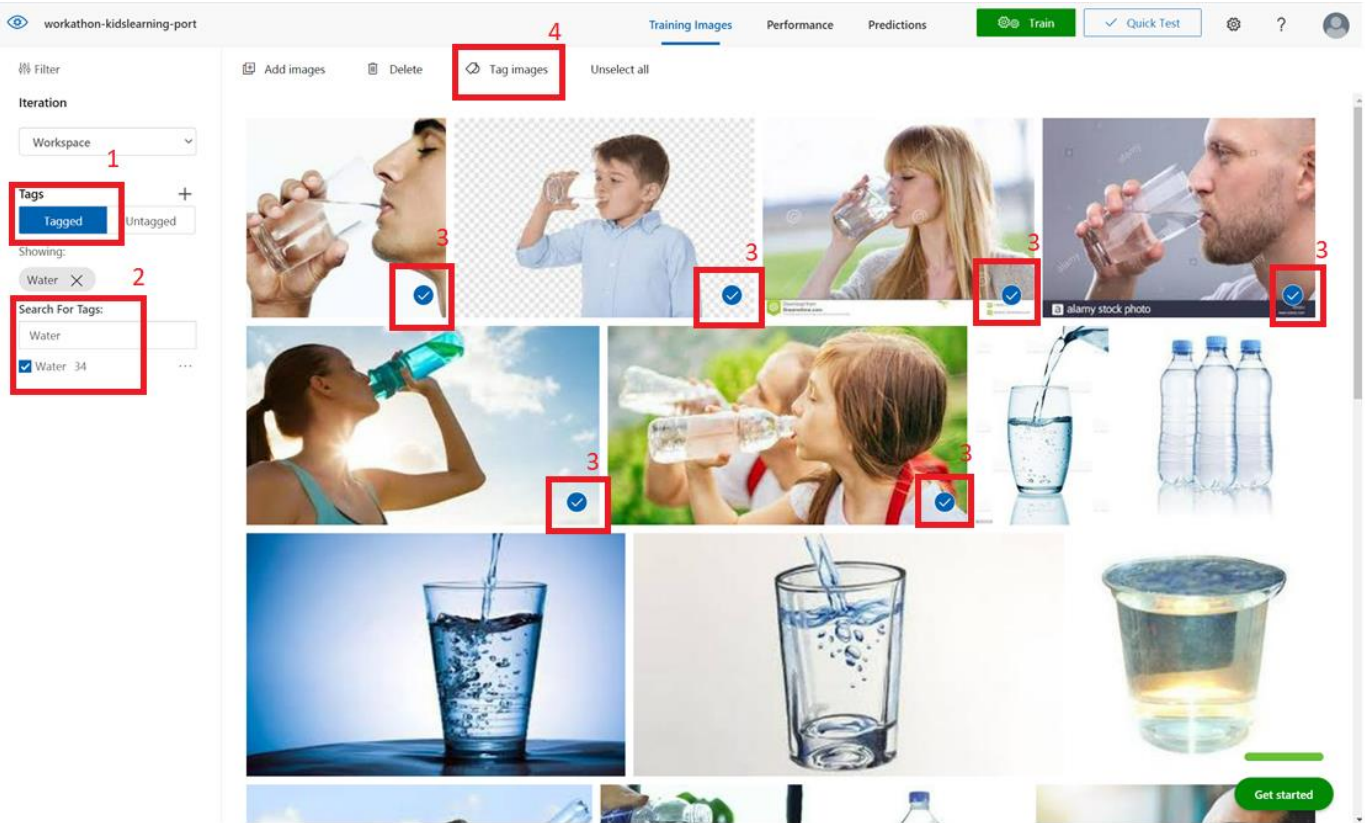
Select images for 1 tag at a time since the tag you will apply (as shown in below image) will apply to all the chosen images.
Repeat the step for other tags as well.

The Activity labels (Drinking, Pouring, Filling) will be done in the next step since they will reuse these images and we don't need to add any other images.





Once you are done adding images and creating tags for Milk, Water & Petrol, you should get a view as shown in the image.



Setting labels for Activity Tags

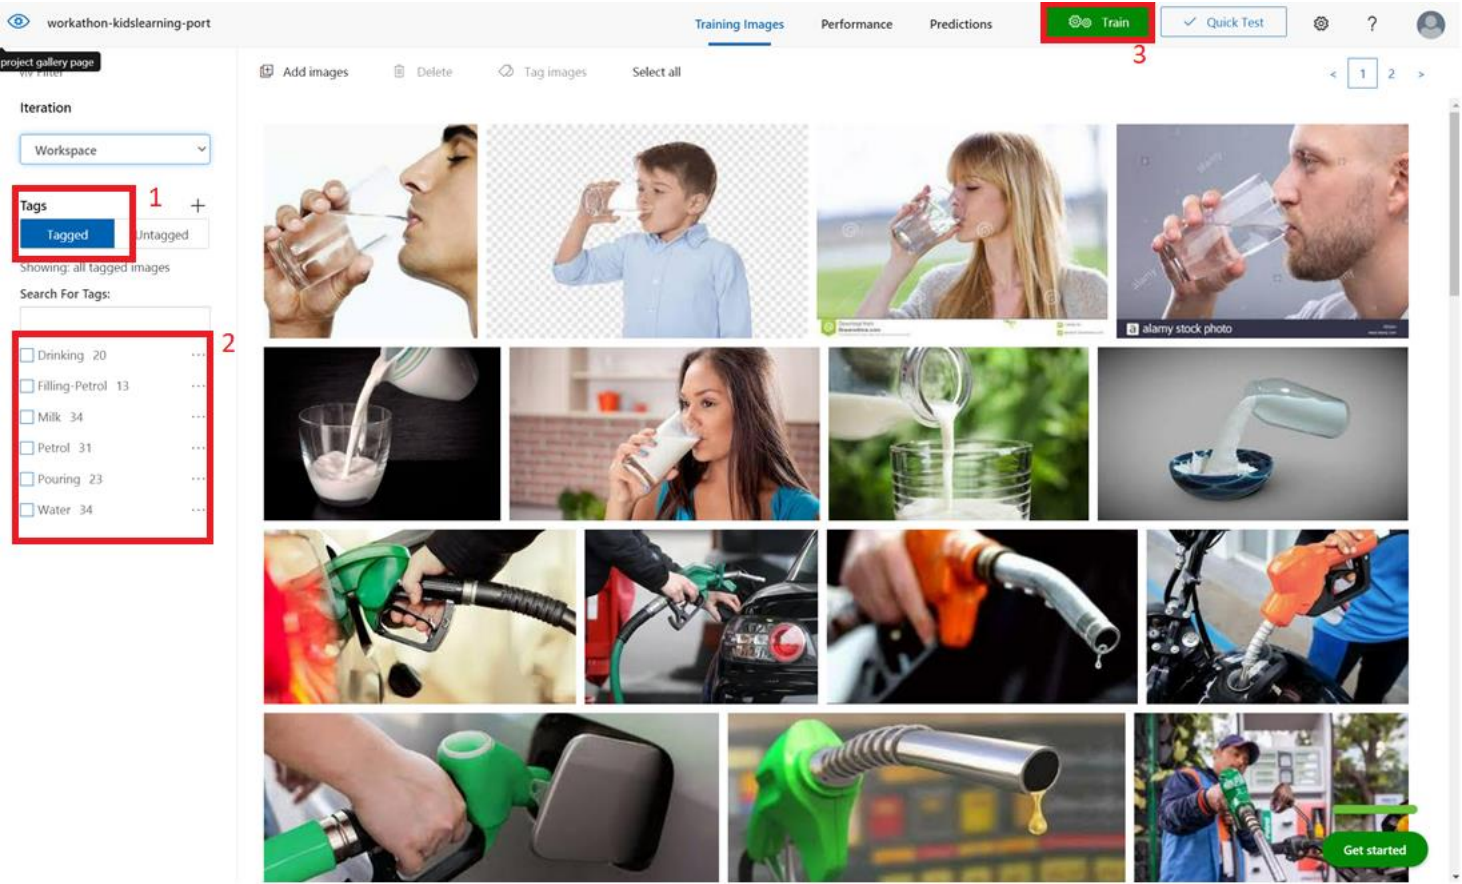
For the 2nd hierarchy of classification – Drinking, Pouring, Filling – follow these steps as highlighted.

For instance, for tagging images from Water label that depict Drinking :

1. Make sure you are on the Tagged images tab (step 1)
2. Filter images to Water so it's easier to find images (step 2)
3. Select all the images that depict drinking (step 3)
4. Click on tag images (step 4)
5. Add drinking tag (as shown in image below)
6. Click Save & Close

Similarly, repeat the above steps for other multi-tagging :

- Pouring in Water images (Filter on water & set Tag as Pouring)
- Pouring in Milk images (Filter on milk & set Tag as Pouring)
- Drinking for Milk images (Filter on milk & set Tag as Drinking)
- Filling in Petrol images (Filter on Petrol & set Tag as Filling)



Training the Model

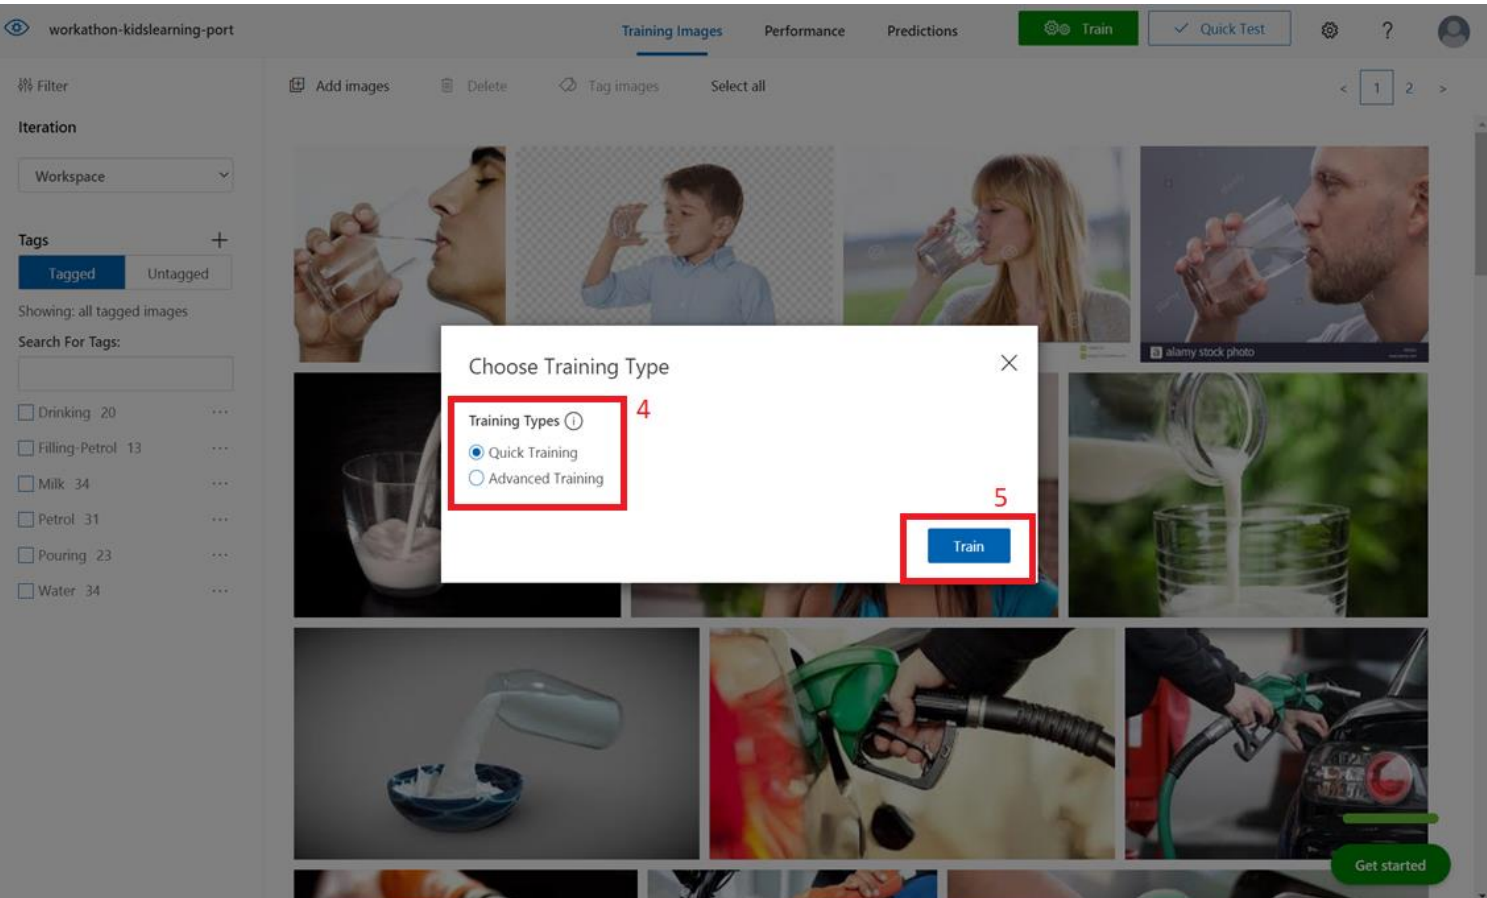
Once you have tagged all images , you'll see the outcome as shown (Step 1,2)

Click the Train the button, to train your custom model (Step 3)

Select Quick Training (Step 4)

We are selecting Quick Training in the interest of time & resources. In a production scenario, you might also want to try out Advanced Training since this will result in higher accuracy.

Click Train.



workathon-kidlearning-port

Training ImagesPerformancePredictionsTrainQuick Test

3

Probability Threshold: 50%

4

Quick Test

Iteration 3

Trained : 3 hours ago with General (compact) domain

Iteration 2

Trained : 4 hours ago with General (compact) domain

Iteration 1

Trained : 4 hours ago with General (compact) domain

Iteration 3

Finished training on 8/13/2021, 10:14:11 AM using General (compact) domain

Iteration id: 90e4611c-eb04-4c22-b142-755a420bf0e5

Classification type: Multilabel (Multiple tags per image)

1

Precision96.0%

Recall80.0%

AP93.6%

2

Performance Per Tag

Tag	Precision	Recall	A.P.	Image count
Water	100.0%	66.7%	93.3%	30
Pouring	100.0%	60.0%	82.5%	24
Petrol	100.0%	100.0%	100.0%	31
Filling-Petrol	100.0%	100.0%	100.0%	13
Drinking	100.0%	100.0%	100.0%	16
Milk	83.3%	71.4%	90.5%	35

Get started

Observe Model Training Outcome & Metrics

Once your model is trained, you will get the outcome as shown. The Precision, Recall & AP metrics measure the performance of your model in different aspects, basis which you can decide how to further tune your model.

Precision is the metrics that tells out of all the images the model labelled as particular tag, how many it labelled correctly.
Eg : In the test set, if it labelled 5 images as water & 4 were actually water, the precision is 80%.
This doesn't mean there were only 5 images of water, there could be more as well that the model didn't catch. For precision, we are just concerned with the ones the model said to be water.

$$Precision = \frac{True\ Positive(TP)}{True\ Positive(TP) + False\ Positive(FP)}$$

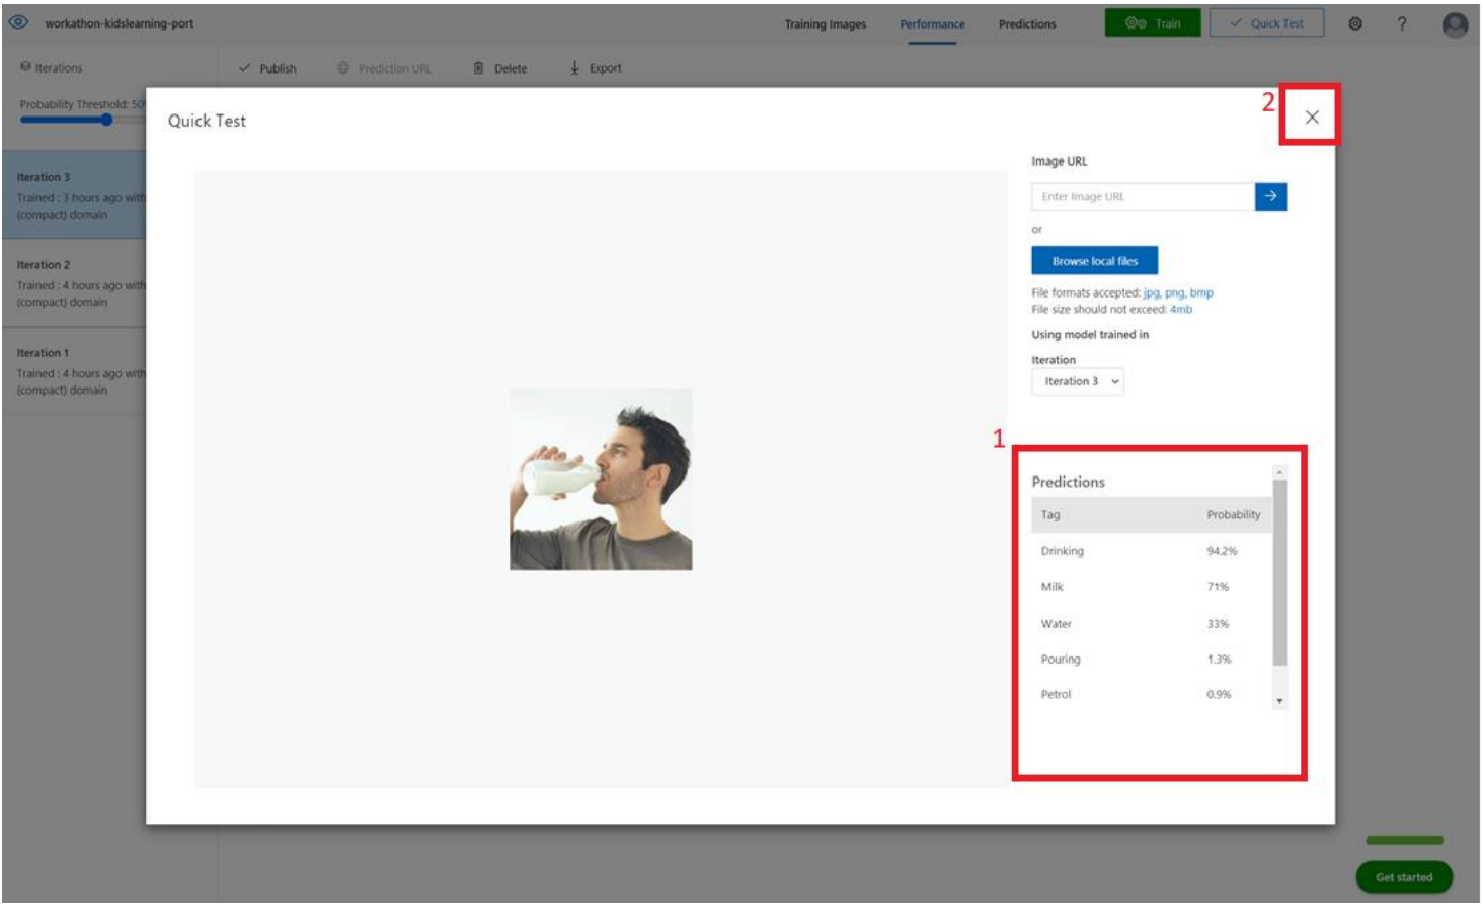
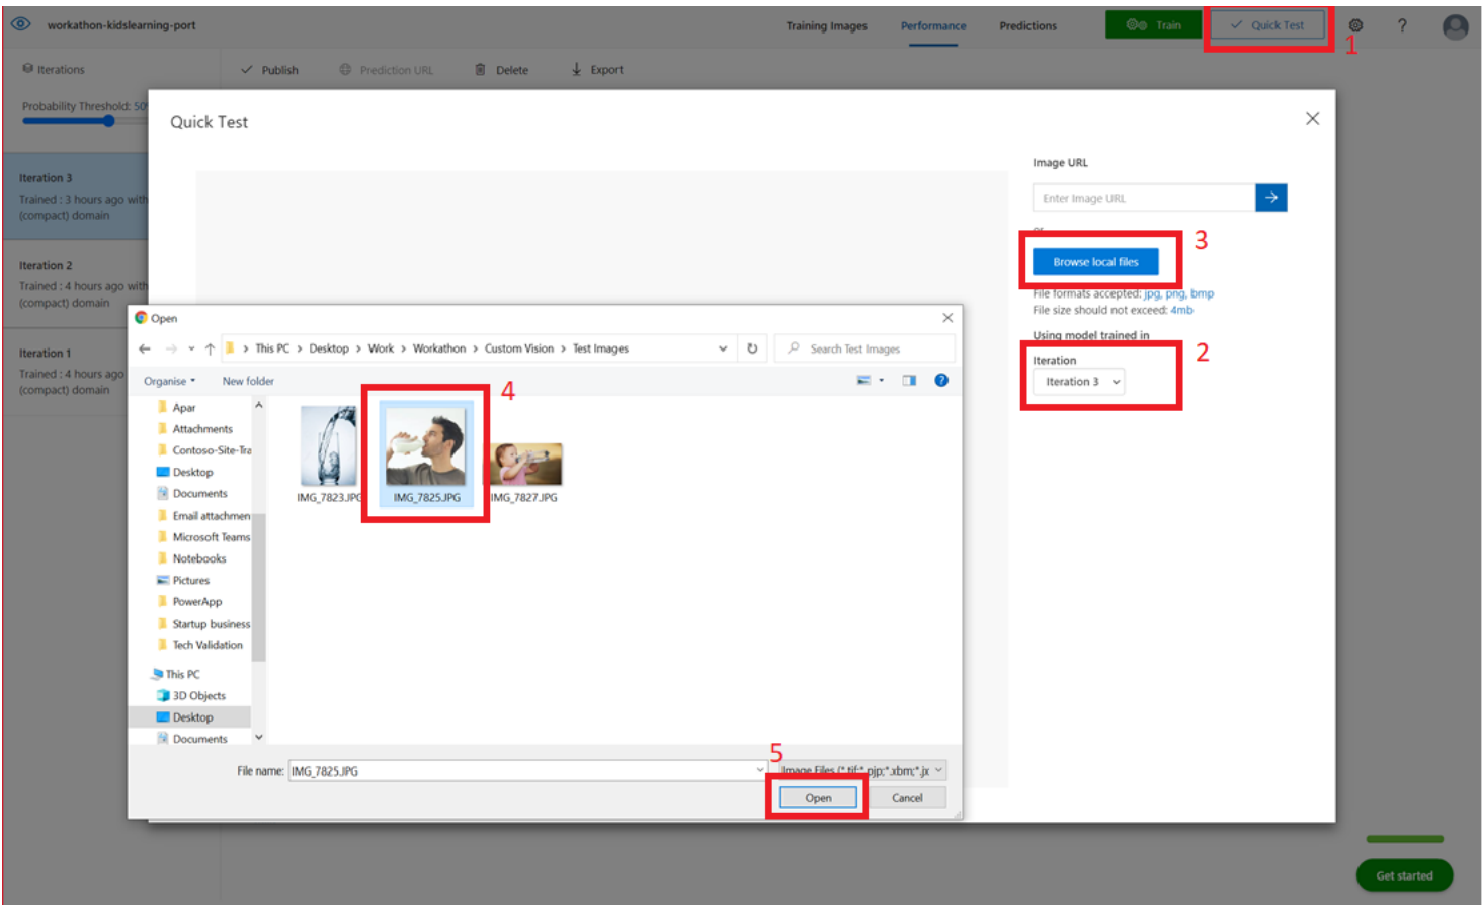
Recall is the metrics that tells out of all the images that belong to a particular tag, how many the model was able to identify.
Eg : In the test set, if there were 8 images of water & it could identify 6, the recall would be 75%.
This doesn't mean the model identified just 6 images as water. It could have identified images belonging to other tags (milk, petrol) as water, however, we are just concerned with water images that it correctly identified.

$$Recall = \frac{True\ Positive(TP)}{True\ Positive(TP) + False\ Negative(FN)}$$

AP is a measure of model performance that summarises the Precision & Recall at different threshold. This is threshold agnostic, hence, changing threshold will not affect AP.

Probability Threshold (step 3) : This is the level of confidence that a prediction needs to have to be considered correct.
Observe the changes in the 3 metrics on changing the threshold slider.
Increasing the threshold will increase Precision but decrease recall. Since now model will label as image as X only if it is super sure, thereby, reducing False Positive & increasing False Negatives.
Similarly, decreasing the threshold will increase Recall but decrease Precision. Since now model will label as image as X even if it is not so sure, thereby, increasing False Positive & decreasing False Negatives.

Now go back & read the definition of Precision & Recall. We're sure you'll have a better understanding now.



Testing the model

Once you have trained the model, you can test it for new images from within the portal itself.

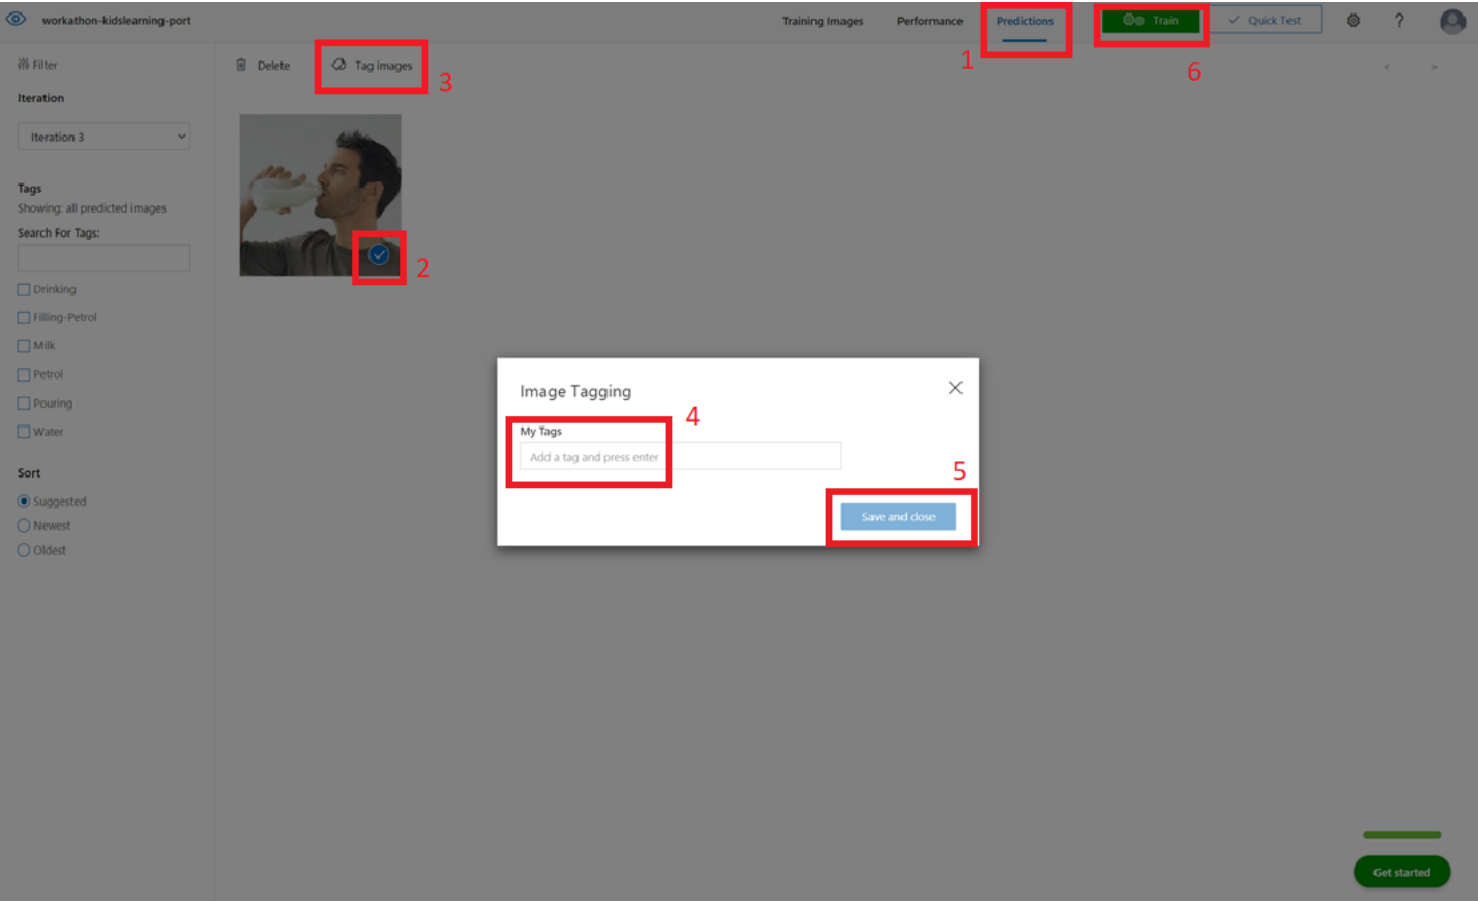
Follow the steps as highlighted:

1. Select quick test tab
2. Select the iteration you want to test (This is important if you have multiple iterations trained. By default, it will pick the latest iteration)
3. Browse local images or provide accessible URLs
4. Click Open.

Once you have selected the image, observe the results & tags predicted and the corresponding probability.

If the results are not as expected or probability confidence is low, you might want to upload more images to the model & retrain.

Any image you test will be available for you to tag & retrain the model. We will see this in the next step.



Adding tested images to model & retraining

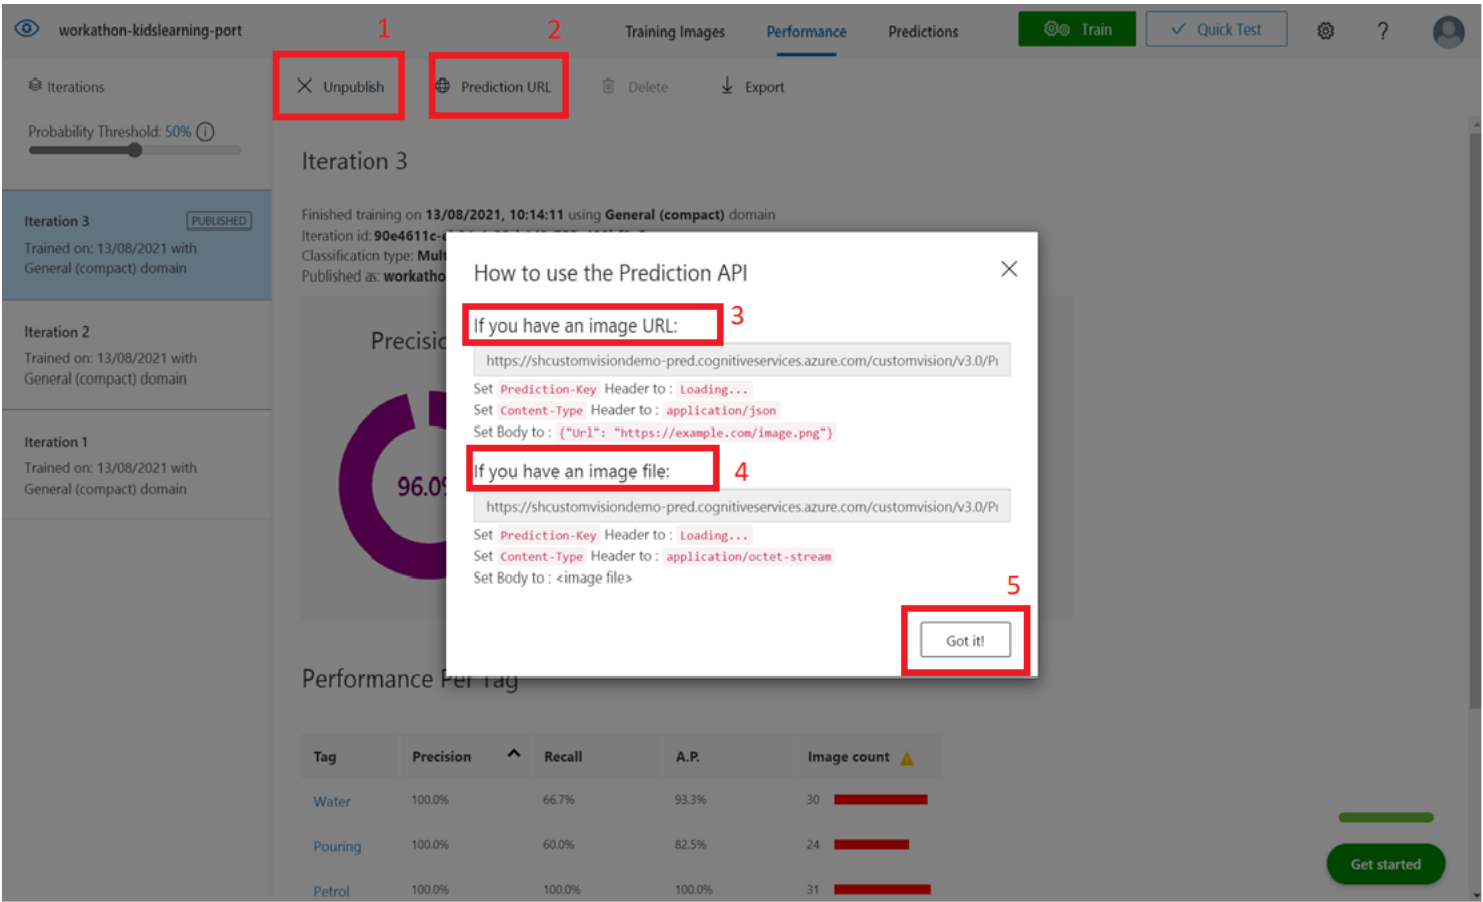
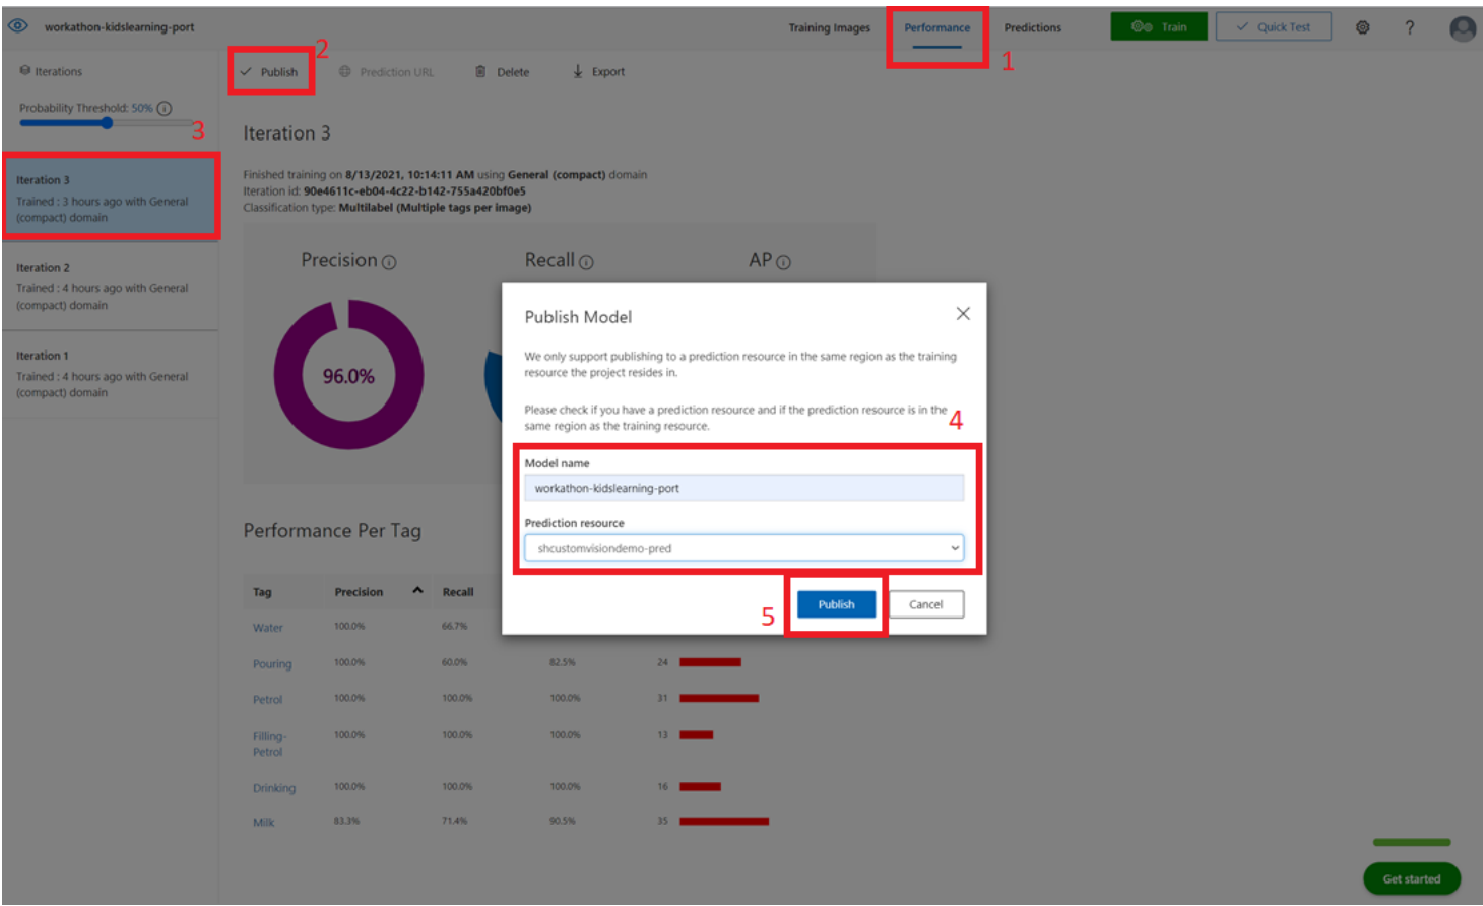
Once you have tested a couple of images in the step above, switch to the Prediction tab as shown in step 1.

Any image you tested in the above step will appear here.

Select all images (step 2) to which you want to apply the same tag and select Tag image (step 3).
Mention the tag you want to apply & click Save and Close (Step 4,5).

Once you have performed steps 2 through 5 for different tags that you want to apply, select Train.

If you want to apply multiple tags to the same image, after applying the first tag, you will have to go to the Training Images section and apply the tags as we applied earlier.



Publishing the Model to Production

When you have done a couple of iterations (Train, test, tune, retrain) and are satisfied with the final iteration, you can publish the model so that it will available for real time predictions.

Follow the highlighted steps:

1. Go to the Performance tab
2. Select Publish
3. Make sure you have selected the right iteration
4. Provide the Model name and Prediction resource (this is the resource we provisioned in the beginning of the workshop : Workathon-customvision-pred)
5. Select Publish

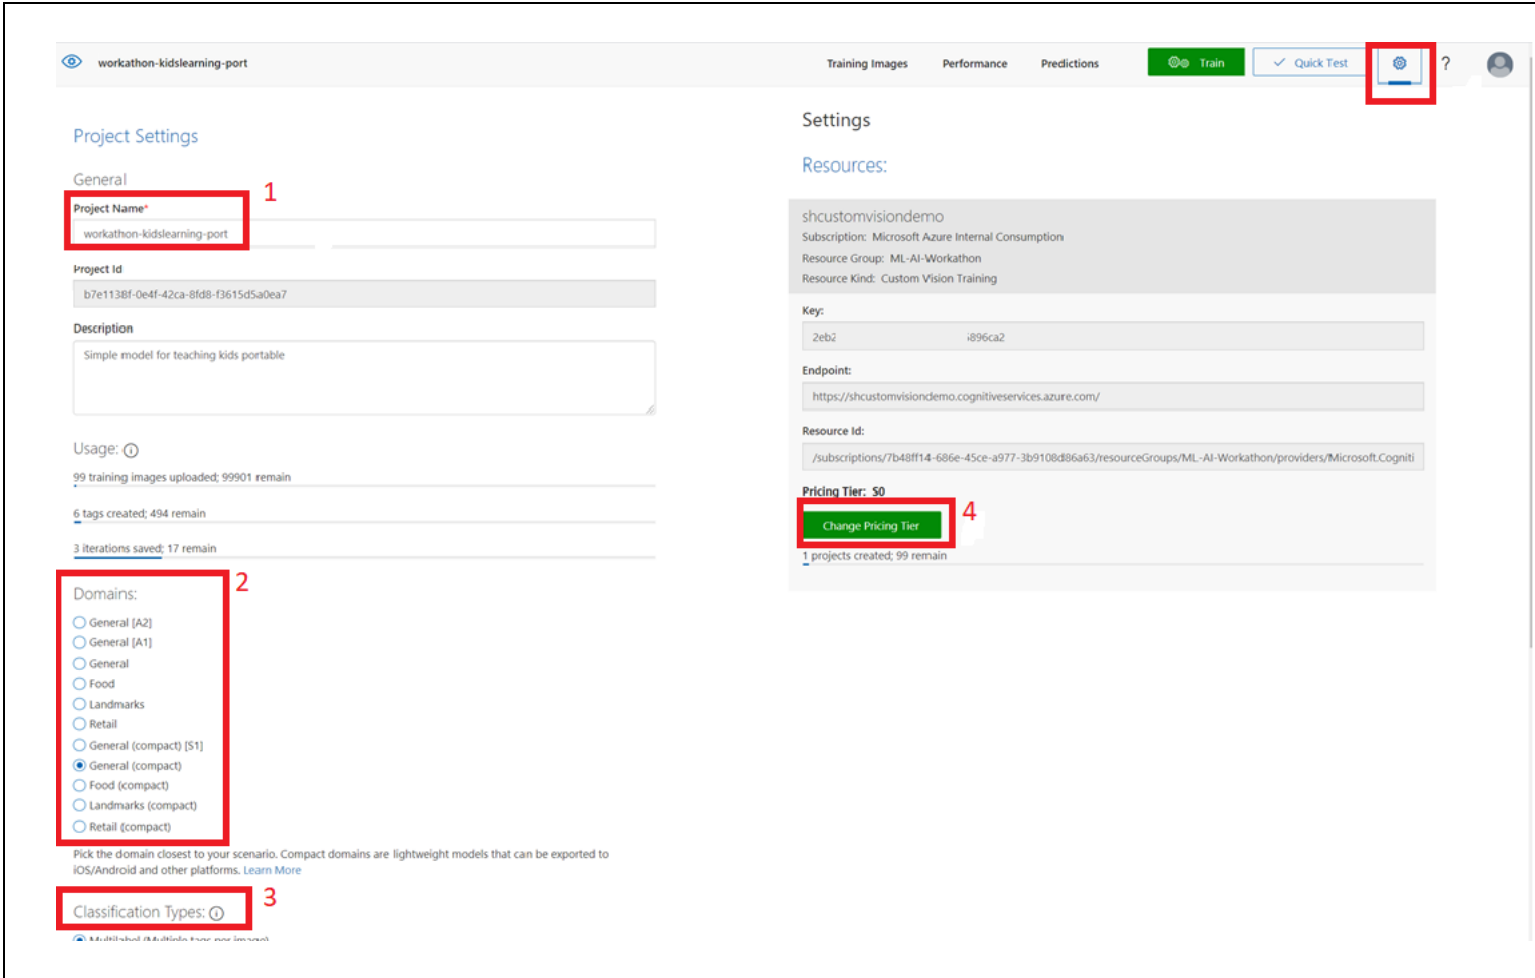
Once published, you will see the option has changed from Publish to Unpublish.

Click the Prediction URL.

Observe you get 2 options to use this in production scenario

1. Provide input as image URL
2. Provide input as image File

We will test this using the REST API in Postman in a later step.



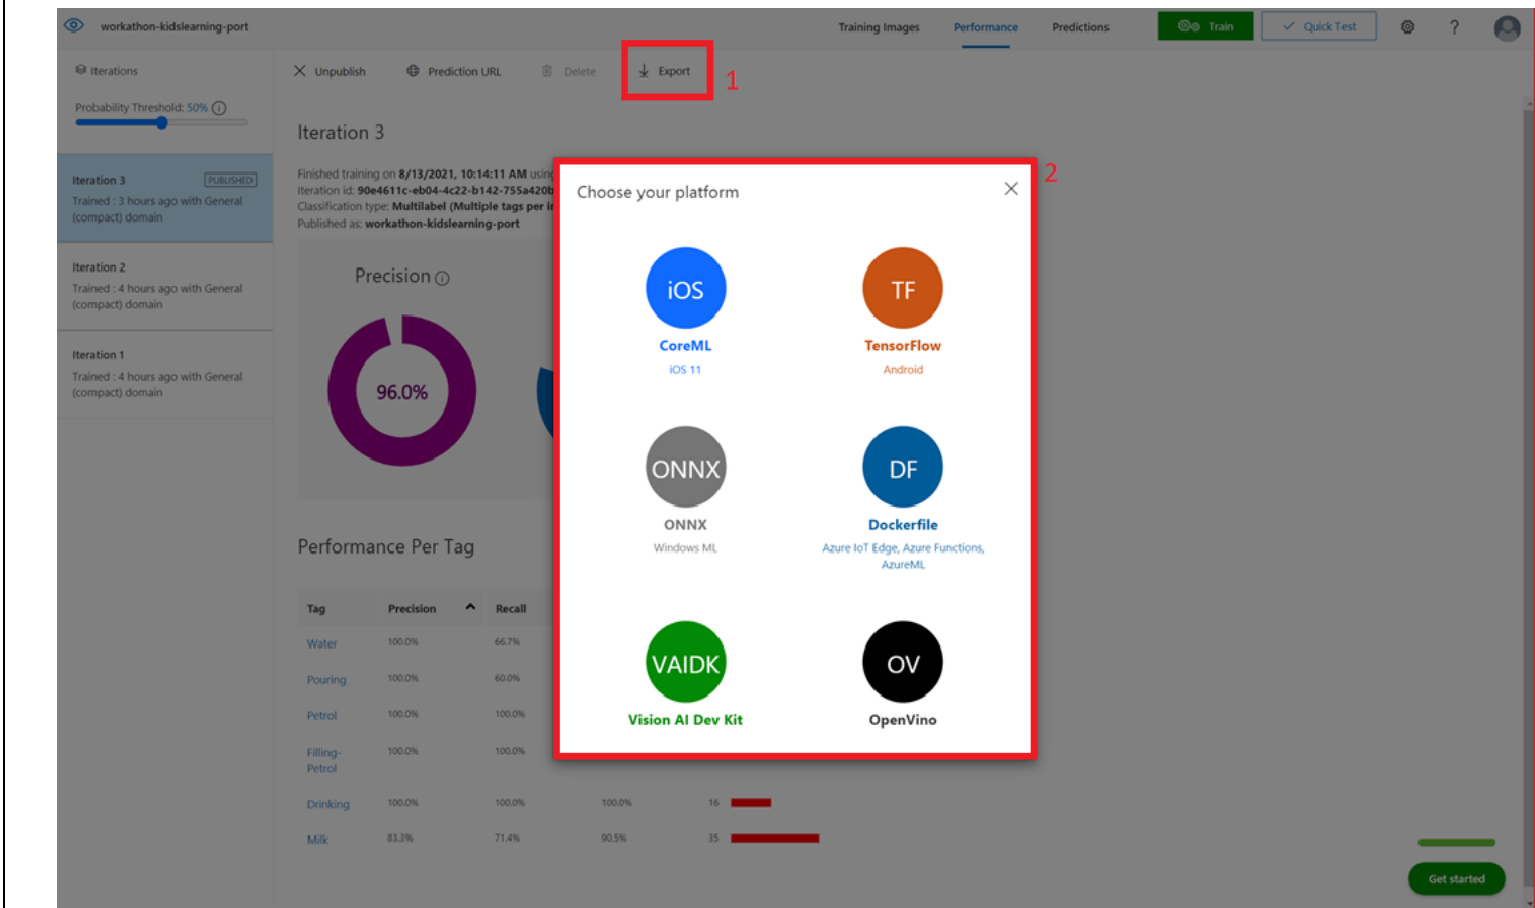
Model & Service Management

Model details can be found in the settings tab.

Some of the details & settings can be changed from this page, such as:

1. Project name & description.
2. Project Domain (You can try different iterations with different domains and see how the performance varies)
3. Classification type (You can switch between multiclass & multilabel)
4. Pricing Tier (in case you started off with free tier and now want to switch to Pay-as-you-Go)

You can also fetch other details such as Project ID, Key, Endpoint etc



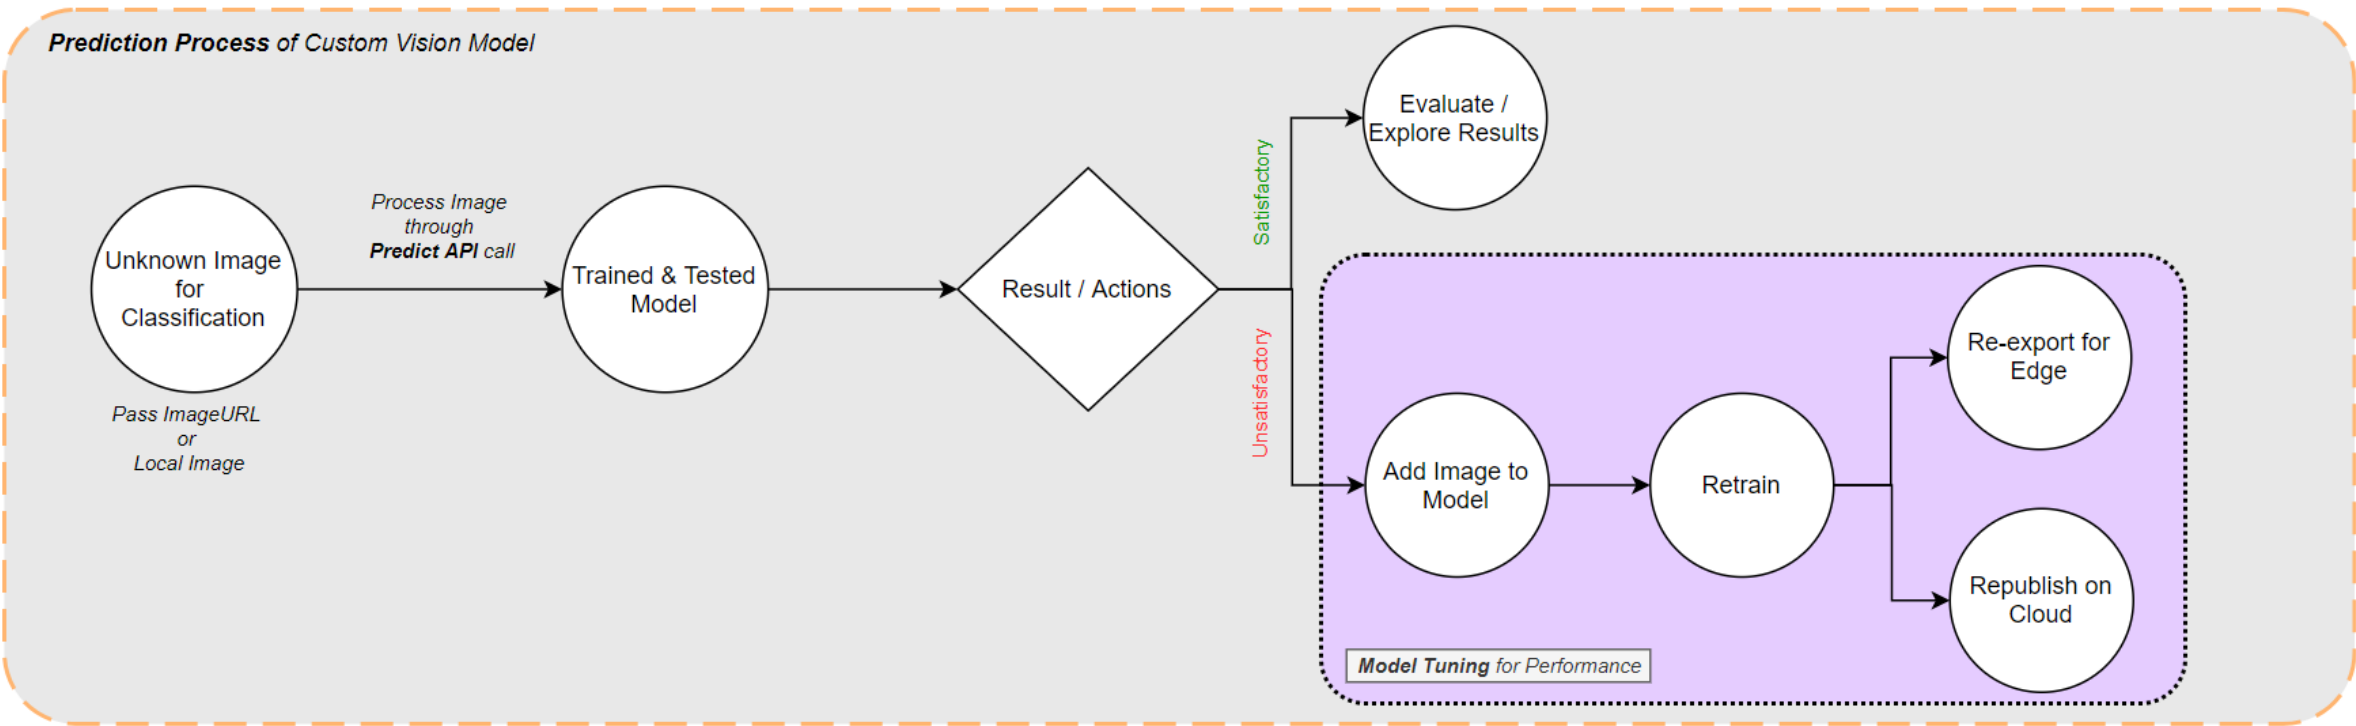
Exporting the model

In the beginning of the workshop, we mentioned we will be selecting the compact model for general domain.

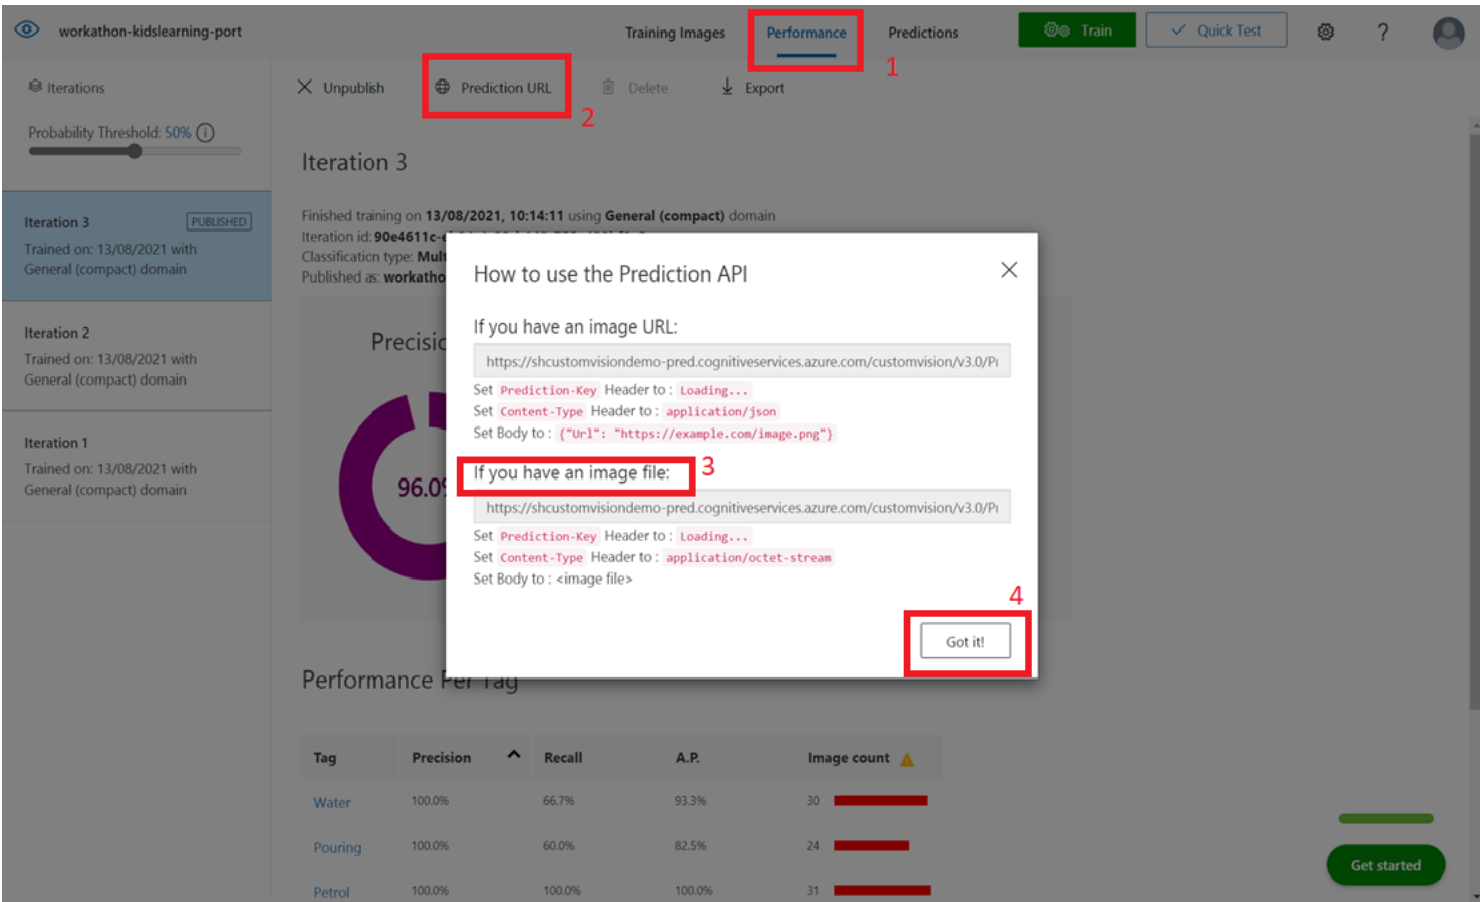
This is where the significance of choosing a compact model comes into play. You can now export the model to run offline in your application or on edge devices.

Select Export (Step 1) & choose the appropriate format for your use case. Artefacts such as model code, required package files etc will be downloaded.

Custom Model Prediction Workflow

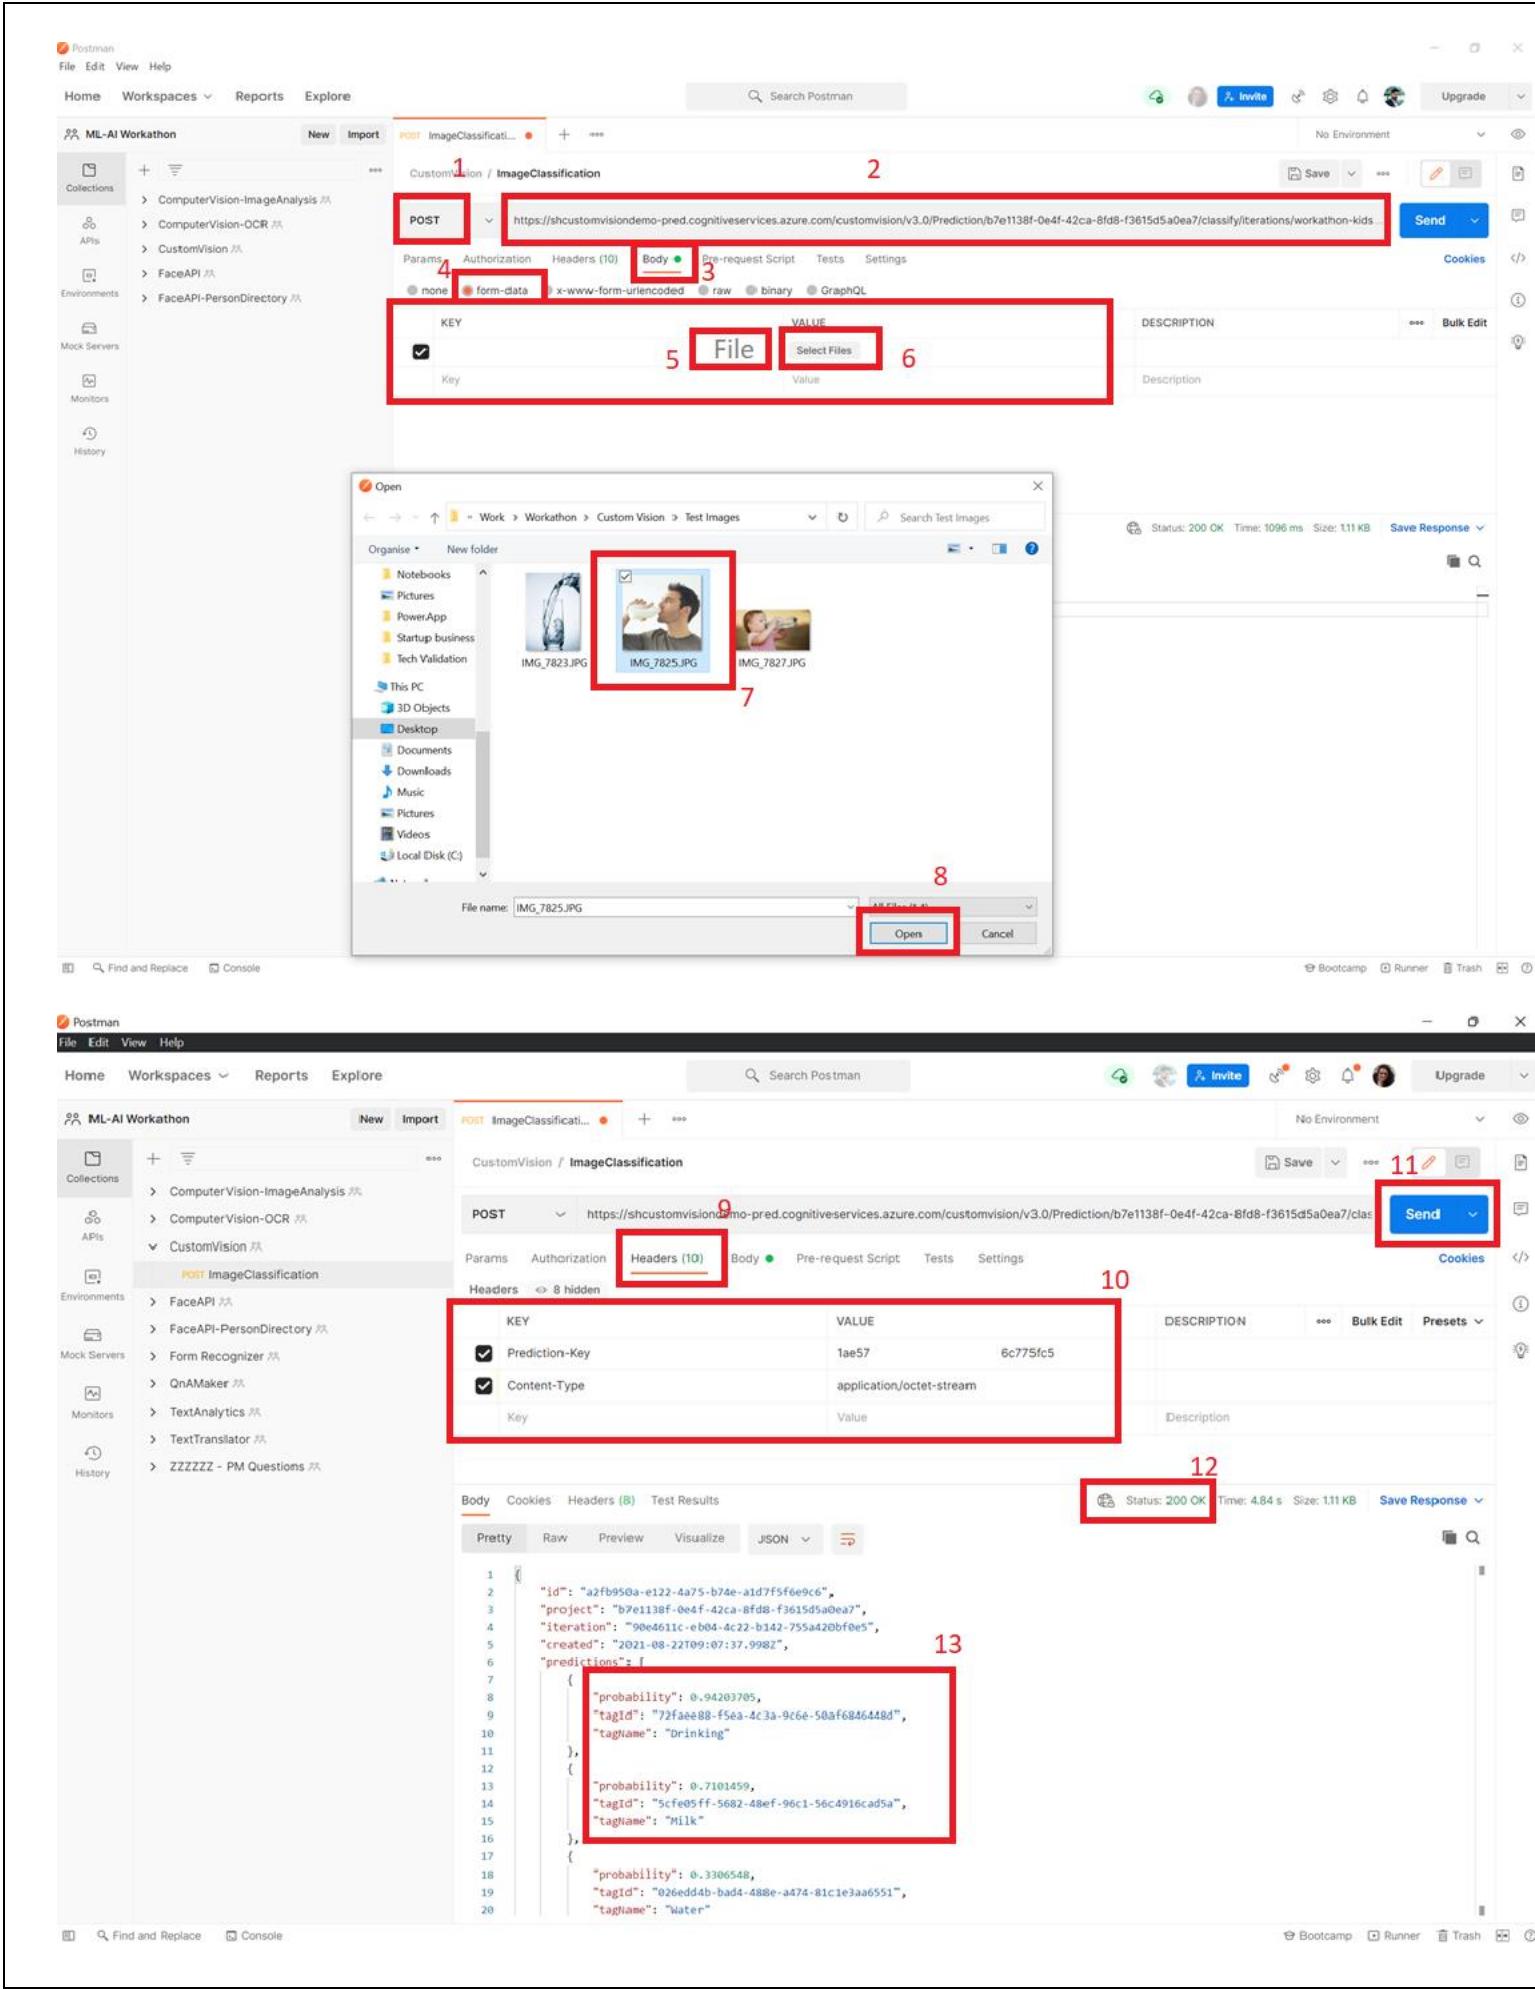


Once you have trained the model, you can leverage it for production scenarios & predict images in real time. Over time, as you predict more and more images, if your model is not performing in the most optimal way, you might also want to leverage those images to further retrain the model. In the below steps, we will predict an image using the REST API Endpoint for the deployed model and add it to the training set.



Classify a new image using your custom model (Using REST APIs in Postman)

From the Performance tab (step 1), select Prediction URL and copy the Endpoint & Headers as shown in step 3.



Fire the Prediction API call using Postman

We will be using Postman to predict a local image. You can use Postman desktop app or web client.

URL : Paste the URL you copied in the above step.

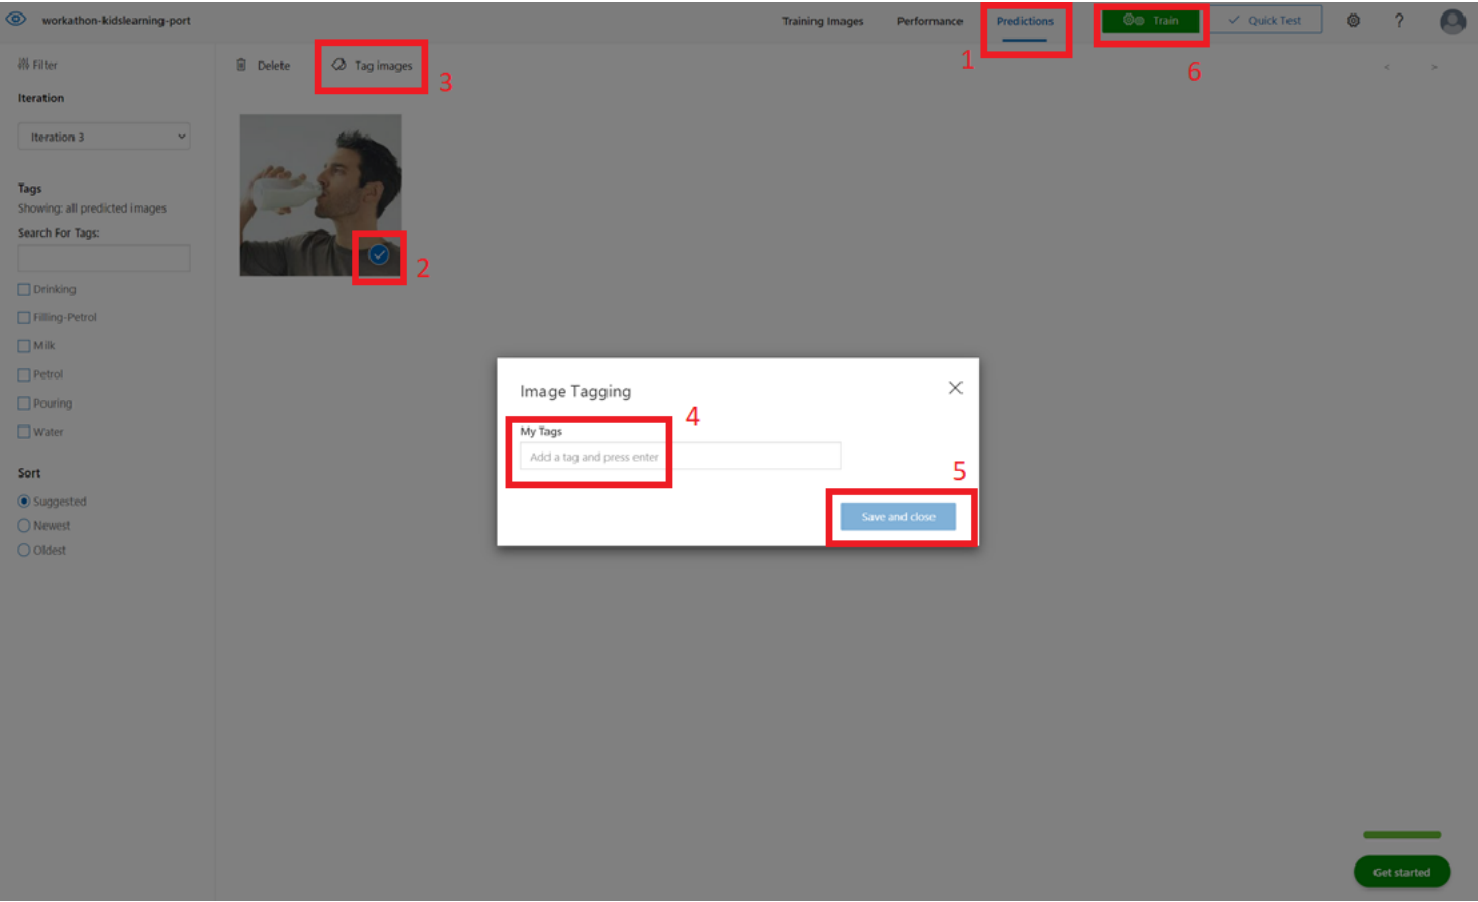
Body :
Select form-data (step 4)
Select file from drop down (step 5)
Browse & select image file you want to predict (step 6-8)

You should also try exploring with different local images & URLs (make sure to change the body & Headers accordingly).

Headers :
Ocp-Apim-Subscription-Key : Paste the key copied in above step
Content-Type : application/octet-stream

Significance of input & output

1. `{{endpoint}}` , `{{key}}` : Values being picked from global variables
2. Ocp-Apim-Subscription-Key : This is the Azure Cognitive service key, that will authenticate the request.
Content-Type : This refers to the input type that you provide in the body, for eg application/octet-stream allows you to enter body in image format. Change the content-type based on input you provide.
3. After you execute the call, observe the status returned, as shown in step 10. This should reflect 200 OK.
4. Observe the predictions section in the JSON output as shown in step 13.



Any image that you make predictions on will be uploaded to your project.

The way we tagged these images and retrained the model earlier, you can follow the steps to repeat the same.

In a production scenario, you can do this activity periodically when you collect a good bunch of images and retrain the model to improve performance.

HOMEWORK

In this workshop, we have explored Image Classification model. We would recommend to you try creating a custom Object Detection Model.

HINT

1. Collect the data set, Create object tags & Tag images
2. Train & test the model, Publish the model
3. Predict on a new image using REST APIs

Additional recommended resources

Service Name	Category	Links
Custom Vision	Programming Language	.Net, Java, Python, Node.js, Go
	Tiers	Free (not for production), Standard
	Pricing	https://azure.microsoft.com/en-us/pricing/details/cognitive-services/custom-vision-service/
	Limits	https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/limits-and-quotas
	Language Support	Not Applicable
	Sample Apps	Sample Skills , Food Recognizer (ONNX) , Video Feed Analysis (OpenCV – IoT) , Rock Paper Scissor
	Regional Availability & Support	https://azure.microsoft.com/en-us/global-infrastructure/services/?products=cognitive-services&regions=all
	SLAs for Cognitive Services	https://azure.microsoft.com/en-in/support/legal/sla/cognitive-services/v1_1/
	Compliance & Certificates	https://azure.microsoft.com/en-us/support/legal/cognitive-services-compliance-and-privacy/
	Cognitive Services Updates	https://azure.microsoft.com/en-us/updates/?product=cognitive-services

Security best practices

1. [Azure Cognitive Services security](#)
2. [Networking](#)
3. [Authentication](#)
4. [Key Management](#)
5. [Data loss prevention](#)
6. [Azure security baseline](#)
7. [Regulatory Compliance controls](#)

Responsible AI being a part of best practices, we encourage you to read [this](#).

[Custom Vision Documentation](#)

[Training API & error link](#)

[Prediction API & error link](#)